

Mining massive Datasets WS 2017/18

Problem Set 1

Rudolf Chrispens, Marvin Klaus, Daniela Schacherer

October 30, 2017

Exercise 01

Given is a cluster of n machines, each having a probability p of failing.

- a) The probability of one machine to not fail is $1 - p$.

The probability of ALL machines not failing is n times $1 - p$ which is $(1 - p)^n$.

The probability of at least one machine failing is the opposite event and thus $1 - (1 - p)^n$.

- b) The probability p_k of exactly k machines failing can be described using the binomial distribution. The binomial distribution describes the discrete probabilities of the number of successes in a sequence of independent experiments. As we have independent machines in the cluster with the number of successes corresponding to a machine failing we can write:

$$p(k|p,n) = \binom{n}{k} p^k (1 - p)^{n-k}$$

p^k is the probability that k machines fail which has to be multiplied to the probability that the other $n - k$ machines do not fail. The binomial coefficient is the combinatoric element and describes in which way k elements can be chosen from n elements.

- c) Zz.: $p_1 + p_2 + \dots + p_n = 1 - (1 - p)^n$

We have $p_1 = p_2 = \dots = p_n = p = \binom{n}{k} p^k (1 - p)^{n-k}$

$$p_1 + p_2 + \dots + p_n = \sum_{k=1}^n \binom{n}{k} p^k (1 - p)^{n-k}$$

We can use the binomial theorem: $\sum_{k=0}^n \binom{n}{k} y^k x^{n-k} = (x+y)^n$ but have to subtract p_0 again

$$\begin{aligned} &= \sum_{k=0}^n \binom{n}{k} p^k (1-p)^{n-k} - \binom{n}{0} p^0 (1-p)^n \\ &= ((1-p) + p)^n - (1-p)^n \\ &= 1^n - (1-p)^n \\ &= 1 - (1-p)^n \end{aligned}$$

Exercise 02

a1) -join() - TRANSFORMATION

Input: otherDataset, [numTasks]

Output: Returns a dataset with "Key/(V1,V2)" pairs.

Code Example `rdd1 = sc.parallelize([("foo", 1), ("bar", 2), ("baz", 3)])`
`rdd2 = sc.parallelize([("foo", 4), ("bar", 5), ("bar", 6)])`
`rdd1.join(rdd2)`

a2) -sort() - TRANSFORMATION - Could not find sort() in reference used sortByKey() instead
-<https://spark.apache.org/docs/2.2.0/rdd-programming-guide.html>

Input: [ascending], [numTasks]

Output: When called on a dataset of (K, V) pairs where K implements Ordered, returns a dataset of (K, V) pairs sorted by keys in ascending or descending order, as specified in the boolean ascending argument.

Code Example: `names = sc.textFile(sys.argv[1])`
`filtered_rows = names.filter(lambdaline : "Count" not inline).map(lambdaline : line.split(", "))`
`filtered_rows.map(lambdan : (str(n[1]), int(n[4]))).sortByKey().collect()`

a2) -groupby() - TRANSFORMATION - Could not find groupby() in reference used groupByKey() instead
-<https://spark.apache.org/docs/2.2.0/rdd-programming-guide.html>

Input: [ascending], [numTasks]

Output: Returns a dataset with "Key/Value" Pairs sorted ascending or descending.

Code Example:

```
lines = spark.read.text(sys.argv[1]).rdd.map(lambda r: r[0])
words = lines.flatMap(lambda x: x.split(' '))
words.reduceByKey(lambda x, y: x + y, 5)
words.groupByKey(5)
```

b1) -NOTE All the tested source code is in *U1Ex2.py*

-INTERSECTION

Input: [RDD]

Output: Returns a RDD with the intersecting elements of two datasets.

Code example:

```
intersectRDD1 = sc.parallelize(range(1, 10))
intersectRDD2 = sc.parallelize(range(5, 15))
intersect = intersectRDD1.intersection(intersectRDD2).collect()
print(intersect)
```

exampleOutput: [8, 9, 5, 6, 7]

b2) -DISTINCT

Input: [numTasks]

Output: Return a new dataset that contains the distinct elements of the source dataset.

example Code:

```
distinctRDD1 = sc.parallelize(range(1, 12))
distinctRDD2 = sc.parallelize(range(8, 20))
distinct = distinctRDD1.union(distinctRDD2).distinct().collect()
print(distinct)
```

exampleOutput: [8, 16, 1, 9, 17, 2, 10, 18, 3, 11, 19, 4, 12, 5, 13, 6, 14, 7, 15]

b3) -UNION

Input: [RDD]

Output: Return a new dataset that contains the union of the elements in the source dataset and the argument.

example Code: `unionRDD1 = sc.parallelize(range(1, 7))
unionRDD2 = sc.parallelize(range(3, 10))
union = unionRDD1.union(unionRDD2).collect()
print(union)`

exampleOutput: `[1, 2, 3, 4, 5, 6, 3, 4, 5, 6, 7, 8, 9]`

b4) -COLLECT

Input: NONE is called as a function on an RDD

Output: Return all the elements of the dataset as an array at the driver program.

example Code: `collection = sc.parallelize([1, 2, 3, 4, 5]).flatMap(lambda x: [x, x, x]).collect()
print(collection)`

exampleOutput: `[1, 1, 1, 2, 2, 2, 3, 3, 3, 4, 4, 4, 5, 5, 5]`

b5) -COUNT

Input: NONE is called as a function on an RDD

Output: Return all the number of elements of the dataset as an array at the driver program.

example Code: `names1RDD = sc.parallelize(["Daniela", "Marvin", "Rudolf", "Kevin", "Jaqueline"])
counts = names1RDD.count()
print(counts)`

exampleOutput: `5`

b6) -FIRST

Input: NONE is called as a function on an RDD

Output: Return all the first element of the dataset as an array at the driver program.

example Code: `names2RDD = sc.parallelize(["Daniela", "Marvin", "Rudolf"])
first = names2RDD.first()
print(first)`

exampleOutput: Daniela

Exercise 03

- a) see comments in 01-03_kmeans.py
- b) see 01-03_kmeans.py

Exercise 04

- a) Broadcast provides a fast way to send data to each desired *node* once. In this context the *node* is one dataset of the data we want to map.
Broadcast is a better solution than to just *join* the data. Once its send to a machine with Broadcast it stays cached on this machine and you can access the Broadcast values on the *nodes*. But be careful, don't modify the Broadcast data.
A *task* is the function we want to execute in the map step.
- b) Accumulators are used to count something up. For example (in the video) we need this to count failures in the application.
If we want to build a custom accumulator we have to implement the type of the accumulator. For example a Vector class to build an accumulator of type Vector.
The accumulator can only be used in the Master (Exception on workers), the `reduce()` function gathers data from all tasks.
- c) The `join()`, `mapValues()` and `reduceByKey()` all results in partitioned RDD's. With partitioning there is less traffic over networks what makes it much faster.
In the pageRank example the links have a partitioner so that links with the same hash are on the same node. To build a custom one you have to implement a class that extends from Partitioner. The class need the variable `numPartitions` and the functions `getPartition()` and `equals()`.

Exercise 05

- Version 1: `[:]` is missing in line 130
In line 130 the variable `centroids` and `newCentroids` would refer to the same instance. In the for-loop `newCentroids` is changed and a new instance with the same values is created with `centroids = newCentroids[:]` in line 157.

- Version 2: `[:]` is missing in line 157

In line 130 `centroids` and `newCentroids` will refer to different instances. In the for-loop `newCentroids` is changed and the variable `centroids` is in line 157 assigned to `newCentroids`, meaning they then refer to the same object. In every further for-loop `newCentroids` will be changed and then assigned to `centroids` although they are already the same instance.

- Version 3: `[:]` is missing in line 130 and line 157

In line 130 the variable `centroids` and `newCentroids` would refer to the same instance. In the for-loop `newCentroids` is changed which changes also `centroids` as they refer to the same object. The same is true for every further for-loop. One of the two variables is therefore needless.