

# Mining Massive Datasets

## Lecture 1

Artur Andrzejak

<http://pvs.ifi.uni-heidelberg.de>



RUPRECHT-KARLS-  
UNIVERSITÄT  
HEIDELBERG



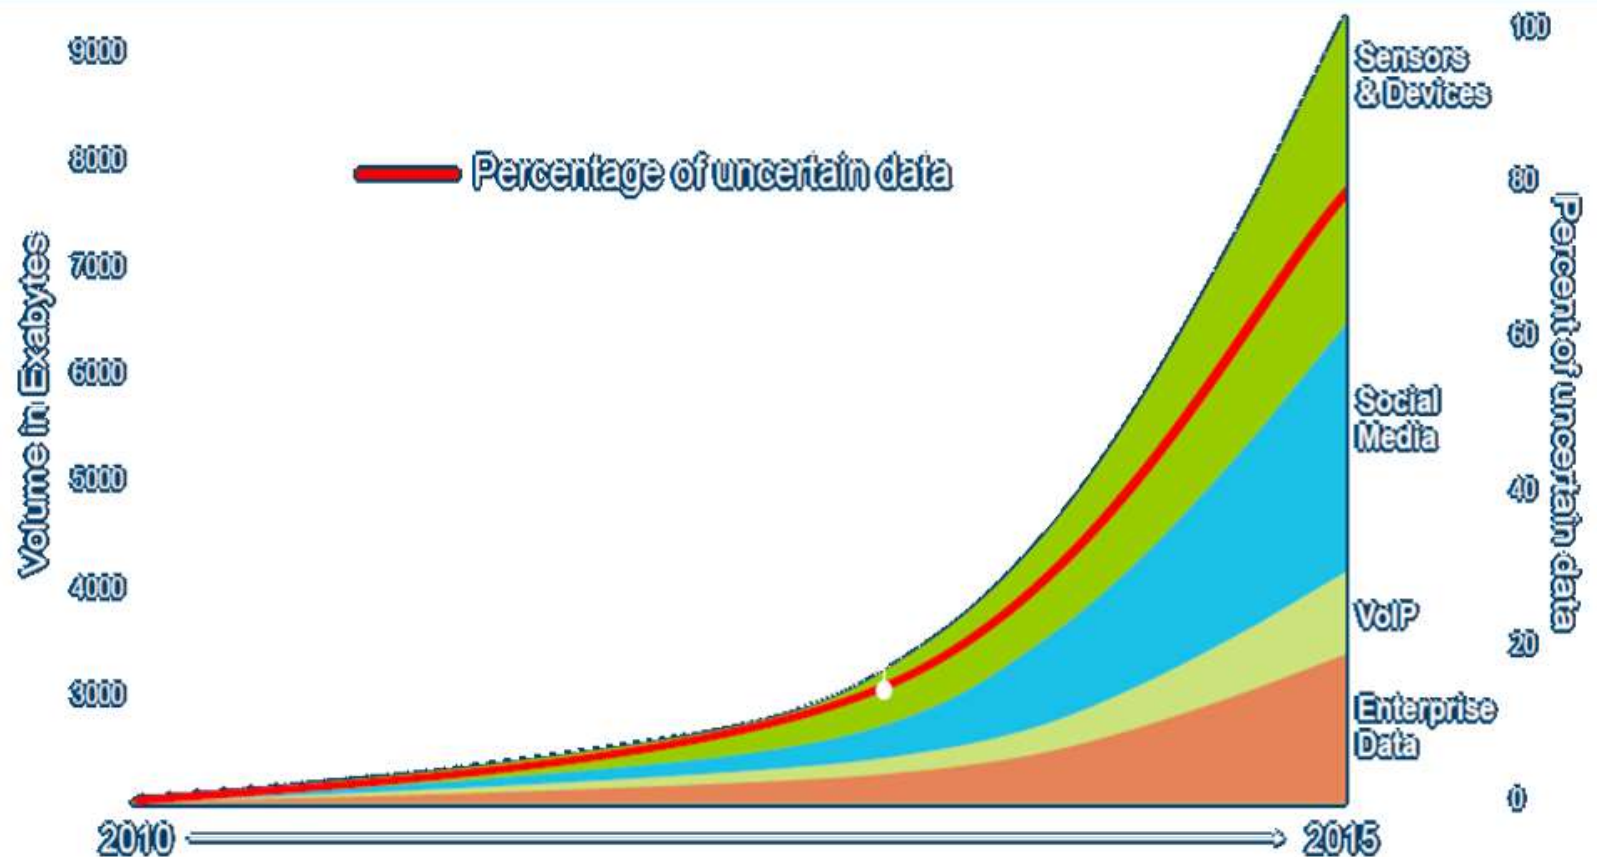
# Note on Slides

A substantial part of these slides come (either verbatim or in a modified form) from the book *Mining of Massive Datasets* by Jure Leskovec, Anand Rajaraman, Jeff Ullman (Stanford University).

For more information, see the website accompanying the book: <http://www.mmds.org>.

# **Motivation and Course Contents**

# Lot of Data, Lot of Computing Power



© 2013 International Business Machines Corporation

## \$5 million vs. \$400

Price of the fastest supercomputer in 1975<sup>1</sup>  
and an iPhone 4 with equal performance

# Data $\neq$ Knowledge

- To extract the knowledge from data it needs to be
    - Stored
    - Managed
    - Processed
    - And Analyzed
- } This course

**Data Mining  $\approx$  Big Data  $\approx$   
Predictive Analytics  $\approx$  Data Science**

# Data Mining Tasks

- **Descriptive methods**


- Find human-interpretable patterns that describe the data
  - **Example:** Clustering

- **Predictive methods**

- Use some variables to predict unknown or future values of other variables
  - **Example:** Recommender systems

# Programming vs. Algorithms

- You need programming to analyze (esp. large) data
  - E.g. Wikipedia text corpus is stored as XML, to extract text you need to write scripts
- In the first part of this course, we will learn **programming** to process and analyze large data
- Later we will use this to implement **algorithms**



# What will we learn?

- **We learn paradigms & tools for programming**
  - Spark & related libs (Shark, GraphX, MLlib, ...)
  - Hadoop MapReduce & related frameworks (Apache Mahout, Pig, ...)
  - ... and a bit of Python
- **We learn to mine different types of data**
  - Data is high dimensional
  - Data is a graph
  - Data is infinite/never-ending
  - Data is labeled



# What will we learn?

- **We will learn to solve real-world problems**
  - Recommender systems
  - Spam detection
  - Duplicate document detection
- **We will learn various algorithmic “tools”**
  - Linear algebra (e.g. SVD)
  - Optimization (stochastic gradient descent)
  - Dynamic programming
  - Hashing (e.g. locality-sensitive hashing, LSH)

# Overview (Superset)

## High dim. data

Locality  
sensitive  
hashing

Clustering

Dimensio-  
nality  
reduction

## Graph data

PageRank,  
SimRank

Community  
Detection

Spam  
Detection

## Infinite data

Filtering  
data  
streams

Web  
advertising

Queries on  
streams

## Machine learning

SVM

Decision  
Trees

Perceptron,  
kNN

## Apps

Recommen  
der systems

Association  
Rules

Duplicate  
document  
detection


Programming in Spark & MapReduce

# Questionnaire

- Who have heard about MapReduce?
- ... about Spark?      25 / 7
- ... about Hive?      3
  
- Who have used Apache Hadoop?      6
  
- Who has programmed in Python?      all
- ... in Java?      25
- ... in Scala?
- ... in C++?

# **Distributed Processing: Motivation**

# Single Node Architecture




Machine Learning, Statistics

“Classical” Data Mining

# Motivation: Google Example

- 20+ billion web pages x 20KB = 400+ TB
- 1 computer reads 30-35 MB/sec from disk
  - ~4 months to read the web
- ~1,000 hard drives to store the web
- Takes even more to **do something useful with the data!**
- **Today, a standard architecture for such problems is emerging:**
  - Cluster of commodity Linux nodes
  - Commodity network (ethernet) to connect them

# Cluster Architecture



Each rack contains 16-64 nodes

In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>

# Large-scale Computing

- Large-scale computing for data mining problems on commodity hardware
- **Challenges:**
  - How do you distribute computation?
  - How can we make it easy to write distributed programs?
  - Dependability - machines fail:
    - One server may stay up 3 years (1,000 days)
    - If you have 1,000 servers, expect to loose 1/day
    - People estimated Google had ~1M machines in 2011
      - 1,000 machines fail every day!



# Storage Infrastructure

- **Problem:**

- If nodes fail, how to store data persistently?

- **Answer:**

- **Distributed File System:**

- Provides global file namespace
    - Google GFS; Hadoop HDFS;

- **Typical usage pattern**

- Huge files (100s of GB to TB)
  - Data is rarely updated in place
  - Reads and appends are common

# Distributed File System

## ■ **Chunk servers**

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

## ■ **Master node**

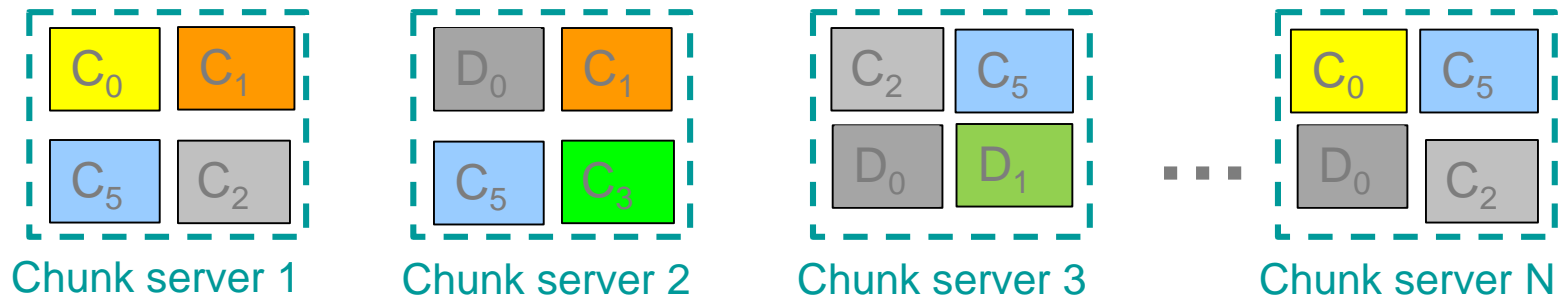
- a.k.a. Name Node in Hadoop's HDFS
- Stores metadata about where files are stored
- Might be replicated

## ■ **Client library for file access**

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

# Distributed File System

- **Reliable distributed file system**
- Data kept in “chunks” spread across machines
- Each chunk **replicated** on different machines
  - Seamless recovery from disk or machine failure



Chunk servers also serve as compute servers

# Implementations

- **Google File System** (GFS)
  - Not available outside Google
  - More info: [Wikipedia](#), [paper](#) (2003)
- **Hadoop Distributed File System** (HDFS)
  - An open-source implementation in Java
  - A de-facto standard for storing large data
  - Download: <http://lucene.apache.org/hadoop/>
- Other distributed file systems
  - [Lustre](#) (for HPC), [Fossil](#) (for HPC), [GPFS](#) (IBM), [Ceph](#) (Red Hat), [XtreemFS](#) (EU Project),...

# What Is Apache Hadoop?

- “The Apache™ Hadoop® project develops open-source software for reliable, scalable, distributed computing.
- ... Modules:
  - **Hadoop Common**
  - **Hadoop Distributed File System (HDFS)**
  - **Hadoop YARN**
  - **Hadoop MapReduce**“
- Other Hadoop-related projects at Apache:
  - [Ambari™](#)
  - [Avro™](#)
  - [Cassandra™](#)
  - [Chukwa™](#)
  - [HBase™](#)
  - [Hive™](#)
  - [Mahout™](#)
  - [Pig™](#)
  - [Spark™](#)
  - [Tez™](#)
  - [ZooKeeper™](#)

# Using HDFS – Various Ways

- Via shell commands, e.g.
  - `bin/hadoop dfs -mkdir <hdfs-dir>`
  - `bin/hadoop dfs -put <local-dir> <hdfs-dir>`
  - `bin/hadoop dfs -get <hdfs-dir> <local-dir>`
- By mounting HDFS like a local file system
  - <http://wiki.apache.org/hadoop/MountableHDFS>
- Programmatically: Java, MapReduce, Spark, ....
  - <https://developer.yahoo.com/hadoop/tutorial/module2.html>
- Watch video “Hadoop Tutorial: Intro to HDFS”
  - <https://www.youtube.com/watch?v=ziqx2hJY8Hg>

# Programming Paradigms

# Large-Scale Data Processing

- Our focus: How to perform large scale computing for data processing & analysis?
- **Challenges:**
  - How to distribute computation?
  - How to make programming less difficult?
  - How to deal with high rate of component failures?



# Efficient Frameworks Needed

- Issues: Copying data costs time / Nodes fails
- Ideas:
  - Store files multiple times for reliability
  - Bring computation close to the data
- **Solutions:**
  - Fault-tolerant storage infrastructure (file system)
    - **Hadoop Distributed File System** (HDFS)
  - Programming models (selection)
    - **Hadoop Map-Reduce**: older, but still most widely used
    - Apache **Spark**: richer set of operators & faster

# A Brief History

MapReduce  
at Google

Hadoop  
at Yahoo!

Paper  
on Spark

2002

2004

2006

2008

2010

2012

2014

MapReduce  
published

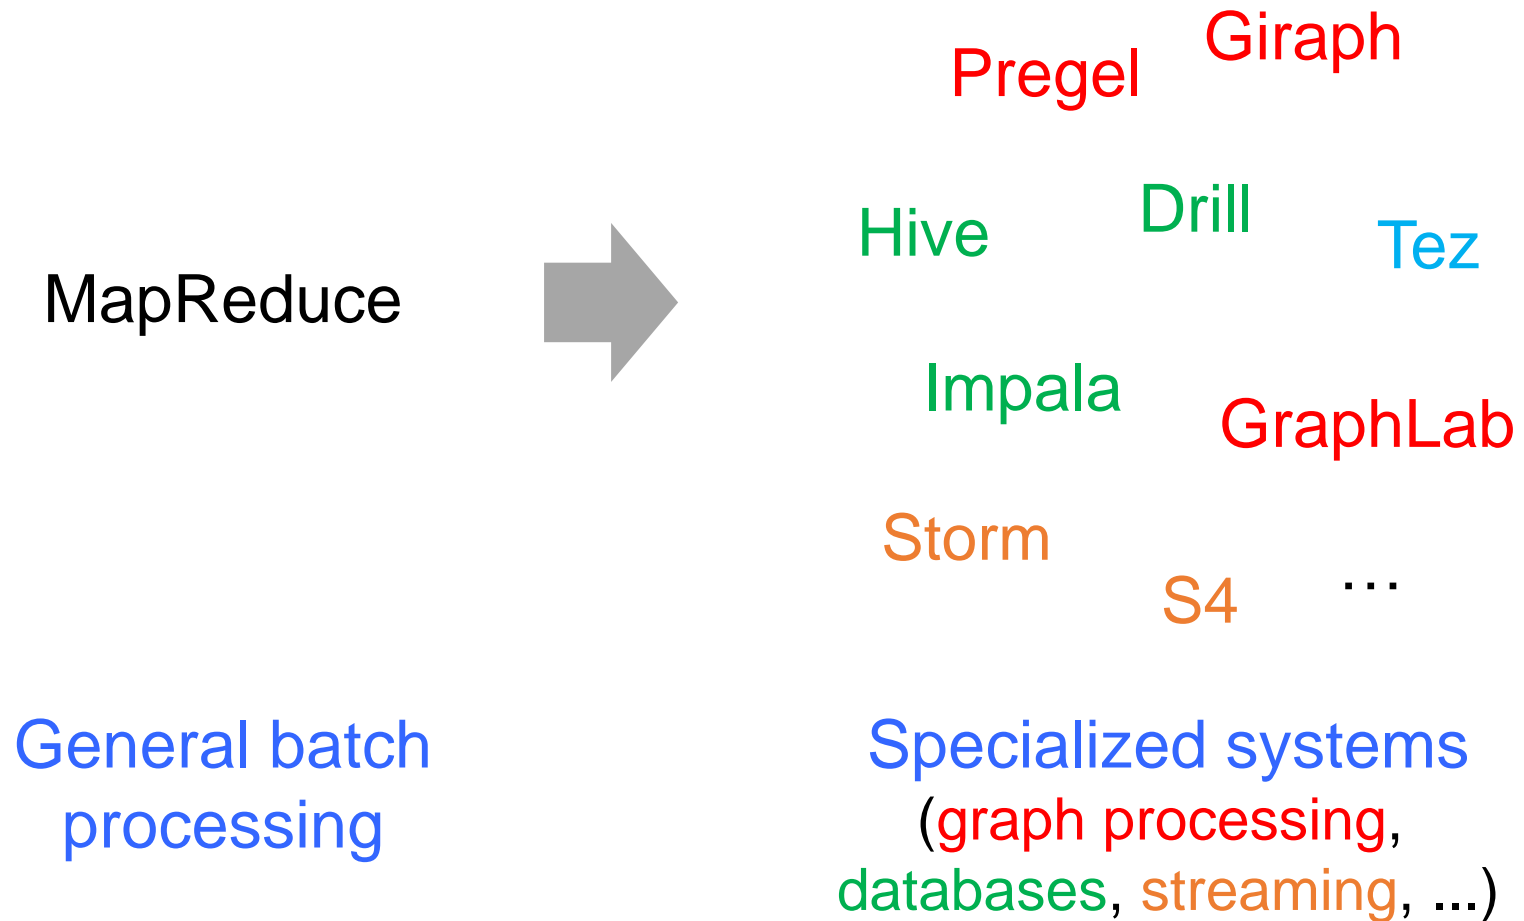
Amazon  
Elastic  
MapReduce  
(2009)

Spark as  
a top-level  
Apache  
project

# MapReduce Problems

- Difficulty of programming
- Performance bottlenecks
  - Esp. for iterative jobs or “multi-MR” jobs
- MapReduce is not an ideal paradigm for large applications (with multiple processing phases)
- => A lot of specialized systems were created as workarounds

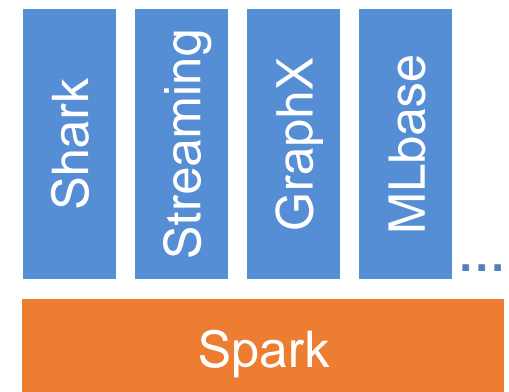
# MapReduce Problems



*From: The State of Spark, and Where We're Going Next*  
**Matei Zaharia**, Spark Summit (2013)

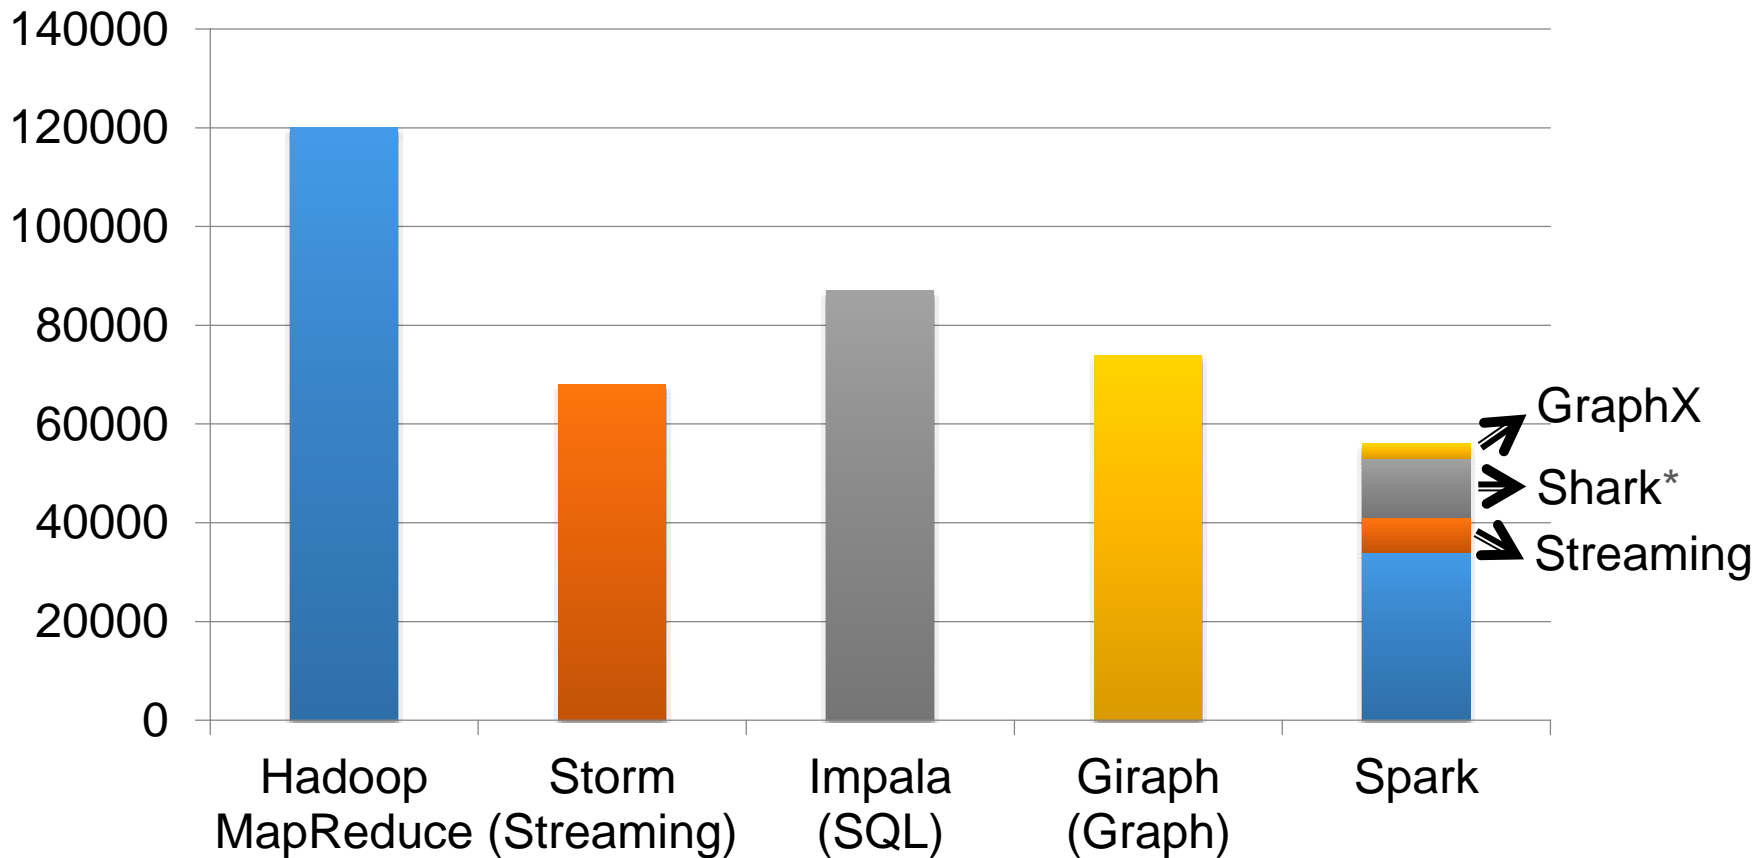
# Spark's Approach

- Instead of specializing, *generalize* MapReduce to support new apps in same engine
- Two changes (general task DAG & data sharing) are enough to express previous models!
- Unification has big benefits
  - For the engine
  - For users



*From: The State of Spark, and Where We're Going Next*  
**Matei Zaharia**, Spark Summit (2013)

# Code Size: Spark vs. Others




non-test, non-example source lines

\* also calls into Hive

*From: The State of Spark, and Where We're Going Next*

**Matei Zaharia**, Spark Summit (2013)

# Performance: Iterative Algorithms



# MapReduce is Being Replaced

*Website of **Mahout**, a library for large-scale machine learning algorithms (as of 10 Oct 2014)*

## 25 April 2014 - Goodbye MapReduce

The Mahout community decided to move its codebase onto modern data processing systems that offer a richer programming model and more efficient execution than Hadoop MapReduce. **Mahout will therefore reject new MapReduce algorithm implementations from now on.**

...

We are building our future implementations on top of a DSL for linear algebraic operations which has been developed over the last months. Programs written in this DSL are automatically optimized and executed in parallel on [Apache Spark](#).



# **Spark Programming: Introduction**

# Virtual Machine with Spark

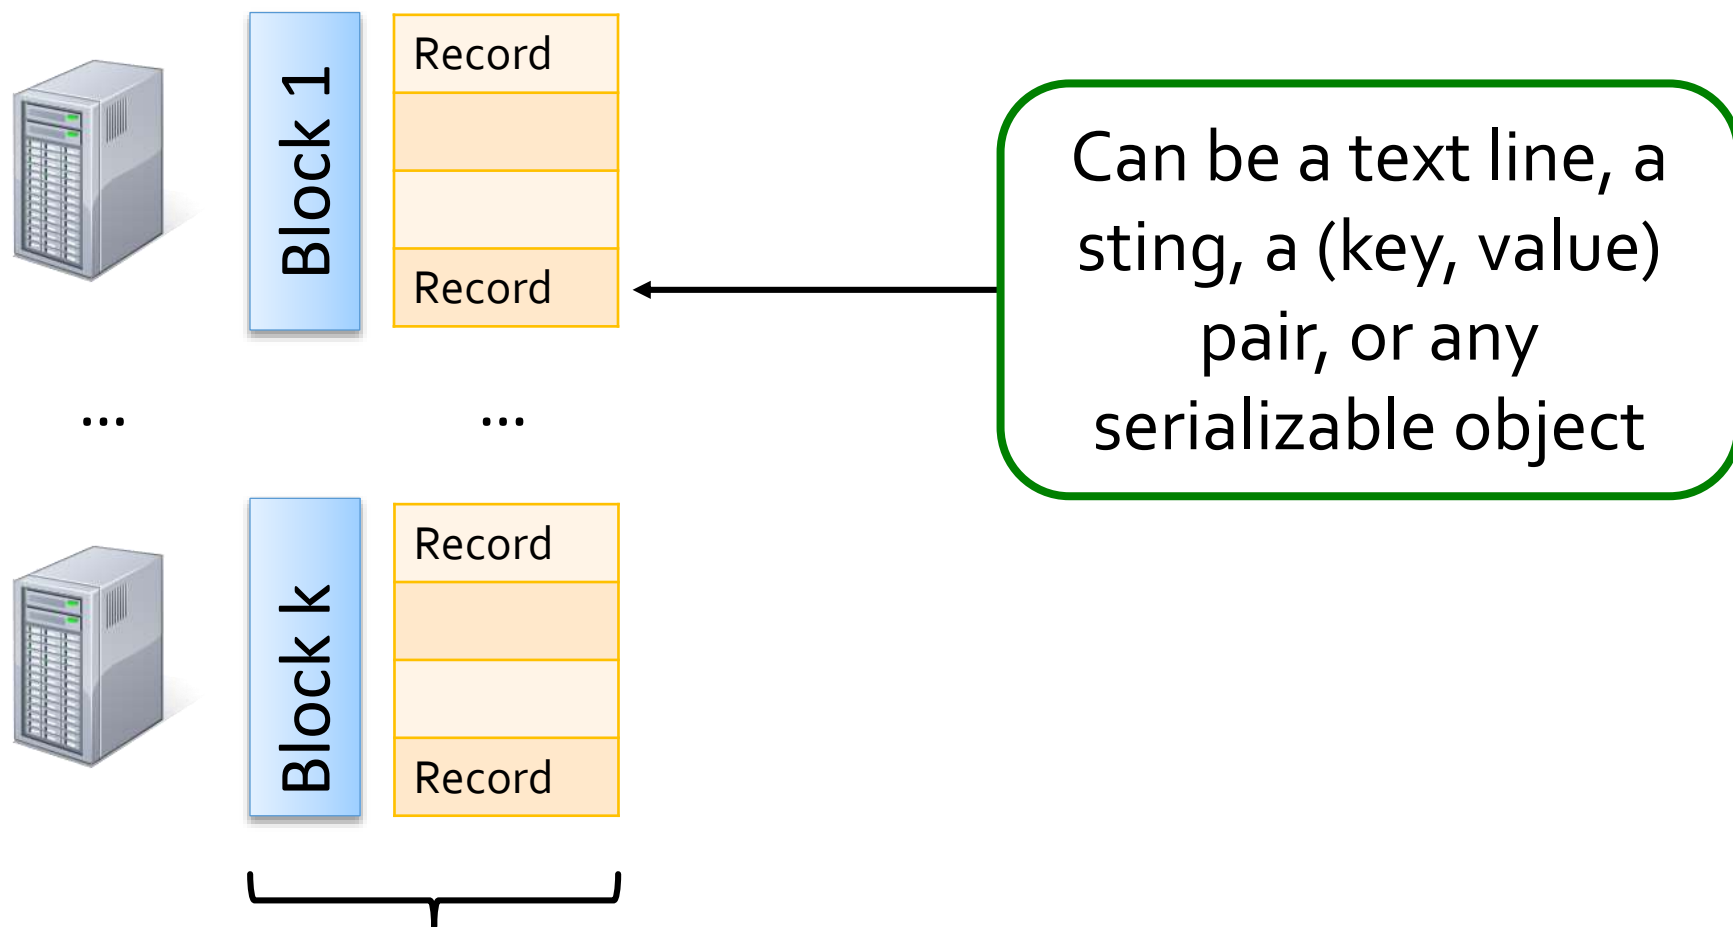
- We have prepared a virtual machine with Spark (under Oracle VirtualBox)
- Available on a USB stick
- Using Spark:
  - Most convenient: included IntelliJ IDEA IDE with a starter project in Python
  - For freaks via a Spark shell:
    - Python: [pyspark](#), Scala: [spark-shell](#)

# Key Idea and Data Structure

- Write programs in terms of **transformations on distributed datasets**
- Main data structure: **resilient distributed dataset (RDD)**
  - Collections of **records** spread across a cluster

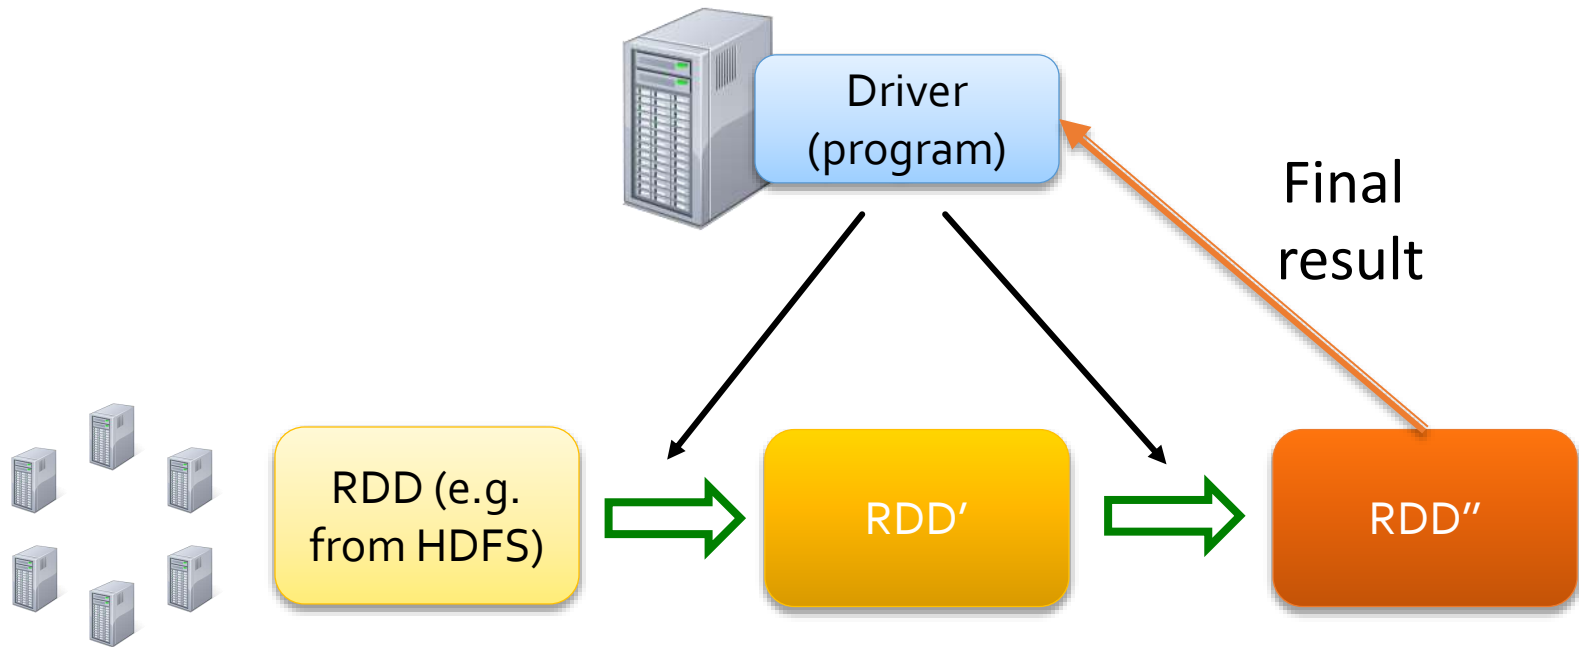


# Resilient Distributed Dataset (RDD)



A RDD; accessible in code via a single “variable”

# Computations with RDDs



Cluster of  
workers

➡ "transformation"

➡ "action"

➡ Code and variables for Ts and As

# Operations

- **Transformations** (e.g. map, filter, groupBy)
  - Lazy operations to build RDDs from other RDDs
- **Actions** (e.g. count, collect, save)
  - Return a result or write it to storage

# Creating RDDs (Python)

**SparkContext** is the “entry point” to Spark functionality, here referenced by variable `sc`

- `# Turn a Python list [1, 2, 3] into an RDD`
- `sc.parallelize([1, 2, 3])`
  
- `# Load text file from local FS (file/dir) or HDFS`
- `sc.textFile("file.txt")`
- `sc.textFile("directory/*.txt")`
- `sc.textFile("hdfs://namenode:9000/path/file")`

# Interlude: Lambdas in Python

- We need to pass code to Ts and As
  - Handle code like variables
- Convenient: **anonymous functions**, or inline functions; in Python: **lambda functions**
- Syntax: 

**lambda** <params> : expression
- Example – computing  $x^2$ 
  - `g = lambda x: x**2`; same as `def g(x): return x**2`
- What is this doing?
  - `lambda x, y: x**2 + y**2 <= 1.0`



# Basic Transformations

- `nums = sc.parallelize([1, 2, 3])`
- `# Pass each element through a function`
- `squares = nums.map(lambda x: x*x) // {1, 4, 9}`
- `# Keep elements passing a predicate`
- `even = squares.filter(lambda x: x % 2 == 0) // {4}`
- `# Map each element to zero or more others`
- `flats = nums.flatMap(lambda x: [x, -x, x*x])`  
`# => {1, -1, 1, 2, -2, 4, 3, -3, 9}`

# Basic Actions

- `nums = sc.parallelize([1, 2, 3])`
- `# Retrieve RDD contents as a local collection`
- `nums.collect() # => [1, 2, 3]`
- `# Return first K elements`
- `nums.take(2) # => [1, 2]`
- `# Count number of elements`
- `nums.count() # => 3`
- `# Merge elements with an associative function`  
`nums.reduce(lambda x, y: x + y) # => 6`
- `# Write elements to a text file`  
`nums.saveAsTextFile("hdfs://file.txt")`

# Working with Key-Value Pairs

- Spark's “distributed reduce” transformations operate on RDDs of **key-value pairs**

- Python:

```
pair = (a, b)
pair[0] # => a
pair[1] # => b
```

- Scala:

```
val pair = (a, b)
pair._1 // => a
pair._2 // => b
```

- Java:

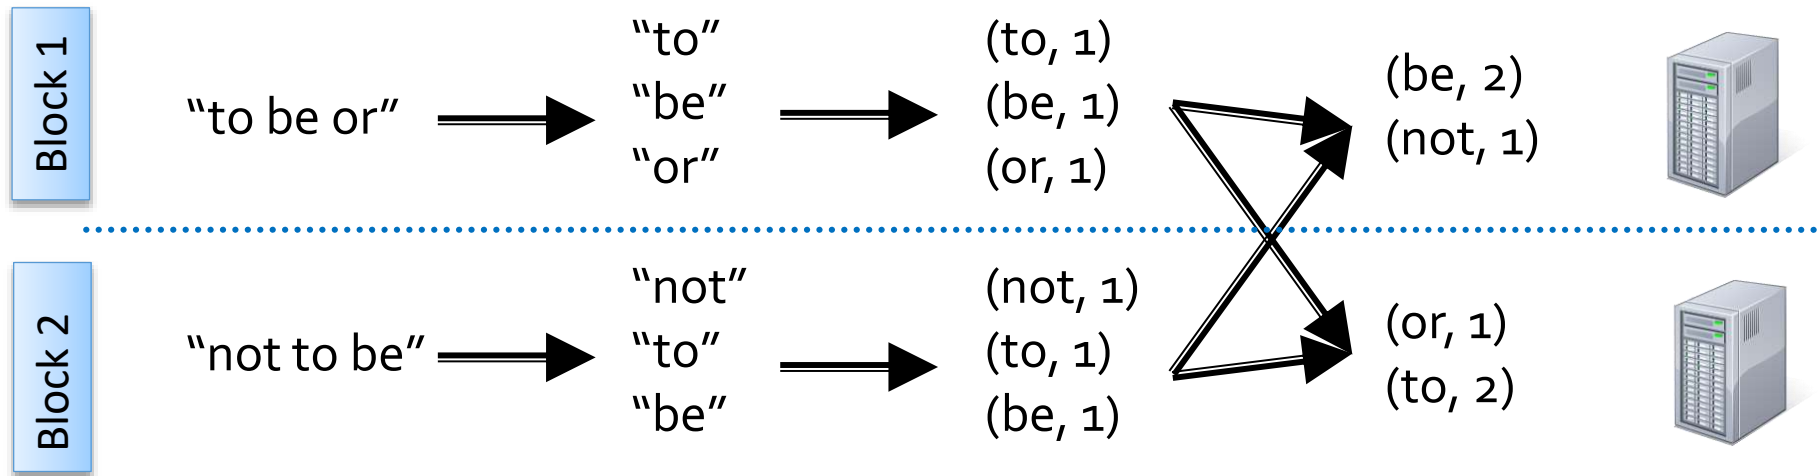
```
Tuple2 pair = new Tuple2(a, b);
pair._1 // => a
pair._2 // => b
```

# Some Key-Value Operations

- `pets = sc.parallelize([("cat", 1), ("dog", 1), ("cat", 2)])`
- `pets.reduceByKey(lambda x, y: x + y)`  
# => {(cat, 3), (dog, 1)}
- `pets.groupByKey()` # => {(cat, [1, 2]), (dog, [1])}
- `pets.sortByKey()` # => {(cat, 1), (cat, 2), (dog, 1)}

# Example: Word Count

- `lines = sc.textFile("hamlet.txt")`
- `counts = lines.flatMap(lambda line: line.split(" "))`  
    `.map(lambda word : (word, 1))`  
    `.reduceByKey(lambda x, y: x + y)`



# Setting the Level of Parallelism

- All the pair RDD operations take an optional second parameter **p** for number of tasks
- `words.reduceByKey(lambda x, y: x + y, 5)`
- `words.groupByKey(5)`

# Using Local Variables

Essentially, piece of code executed by a transformation or action

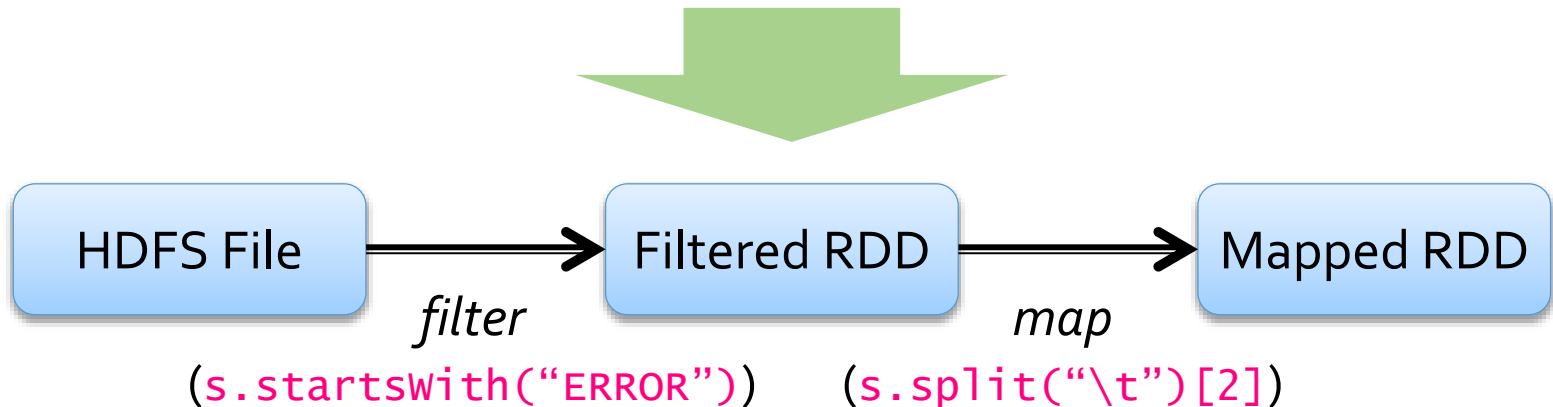
- Any external variables you use in a **closure** will automatically be shipped to the cluster:
- `query = sys.stdin.readline()`  
`pages.filter(lambda x: query in x).count()`
- Some caveats:
  - Each task gets a new copy (no updates sent back)
  - Variable must be Serializable / Pickle-able
  - Don't use fields of an outer object (ships all of it!)

There are also shared variables:  
**broadcasts, accumulators**

# Fault Recovery

RDDs track *lineage* information that can be used to efficiently recompute lost data

Ex: `msgs = textFile.filter(lambda s: s.startsWith("ERROR"))  
.map(lambda s: s.split("\t")[2])`





# Other RDD Operators

- map
- filter
- groupBy
- sort
- union
- join
- leftOuterJoin
- rightOuterJoin
- reduce
- count
- fold
- reduceByKey
- groupByKey
- cogroup
- cross
- zip
- sample
- take
- first
- partitionBy
- mapWith
- pipe
- save
- ...

More details: [spark-project.org/docs/latest/](http://spark-project.org/docs/latest/)

# Recommended Videos on Spark

- Introduction tutorials on Spark

- Parallel programming with Spark Presented by Matei Zaharia UC Berkeley AmpLab 2013

- <https://www.youtube.com/watch?v=e-56inQL5hQ&t=30s>

- Parallel Programming with Spark (Part 1 & 2) by Matei Zaharia (2012)

- <https://www.youtube.com/watch?v=7k4yDKBYOcw>

- Coursera

- Big Data Analysis with Scala and Spark

- <https://www.coursera.org/learn/scala-spark-big-data>

- Enroll -> Audit, then for free!

# Reading Materials on Spark

## ■ Free Materials:

### ■ Spark Programming Guide

- <https://spark.apache.org/docs/latest/rdd-programming-guide.html>

### ■ Apache Spark Tutorial: ML with PySpark

- <https://goo.gl/u4RjeB>

### ■ Cheat Sheet PySpark-RDD Basics, <https://goo.gl/UF5zVr>

### ■ Jacek Laskowski, Mastering Apache Spark 2, GitBook.com, <https://goo.gl/yFYRYm>

## ■ Books

### ■ Matthew Rathbone: 10+ Great Books for Apache Spark

- <https://blog.matthewrathbone.com/2017/01/13/spark-books.html>

# More Resources

- Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung: **The Google File System**
  - <http://labs.google.com/papers/gfs.html>
- **Hadoop Wiki**
  - Introduction: <http://wiki.apache.org/lucene-hadoop/>
  - Getting Started: <http://wiki.apache.org/lucene-hadoop/GettingStartedWithHadoop>
  - Map/Reduce Overview
    - <http://wiki.apache.org/lucene-hadoop/HadoopMapReduce>
    - <http://wiki.apache.org/lucene-hadoop/HadoopMapRedClasses>

**Thank you.**

Questions?