

# Mining Massive Datasets

## Lecture 5

Artur Andrzejak

<http://pvs.ifi.uni-heidelberg.de>



RUPRECHT-KARLS-  
UNIVERSITÄT  
HEIDELBERG



# Note on Slides

A substantial part of these slides come (either verbatim or in a modified form) from the book *Mining of Massive Datasets* by Jure Leskovec, Anand Rajaraman, Jeff Ullman (Stanford University).

For more information, see the website accompanying the book: <http://www.mmds.org>.

# Current Topic

## High dim. data

Locality  
sensitive  
hashing

Clustering

Dimensio-  
nality  
reduction

## Graph data

PageRank,  
SimRank

Community  
Detection

Spam  
Detection

## Infinite data

Filtering  
data  
streams

Web  
advertising

Queries on  
streams

## Machine learning

SVM

Decision  
Trees

Perceptron,  
kNN

## Apps

Recommender  
systems


Association  
Rules

Duplicate  
document  
detection


Programming in Spark & MapReduce

# **Recommender Systems: Content-based Systems & Collaborative Filtering**

# Example: Recommender Systems



Faculty of  
Math & CS



Faculty of  
Physics

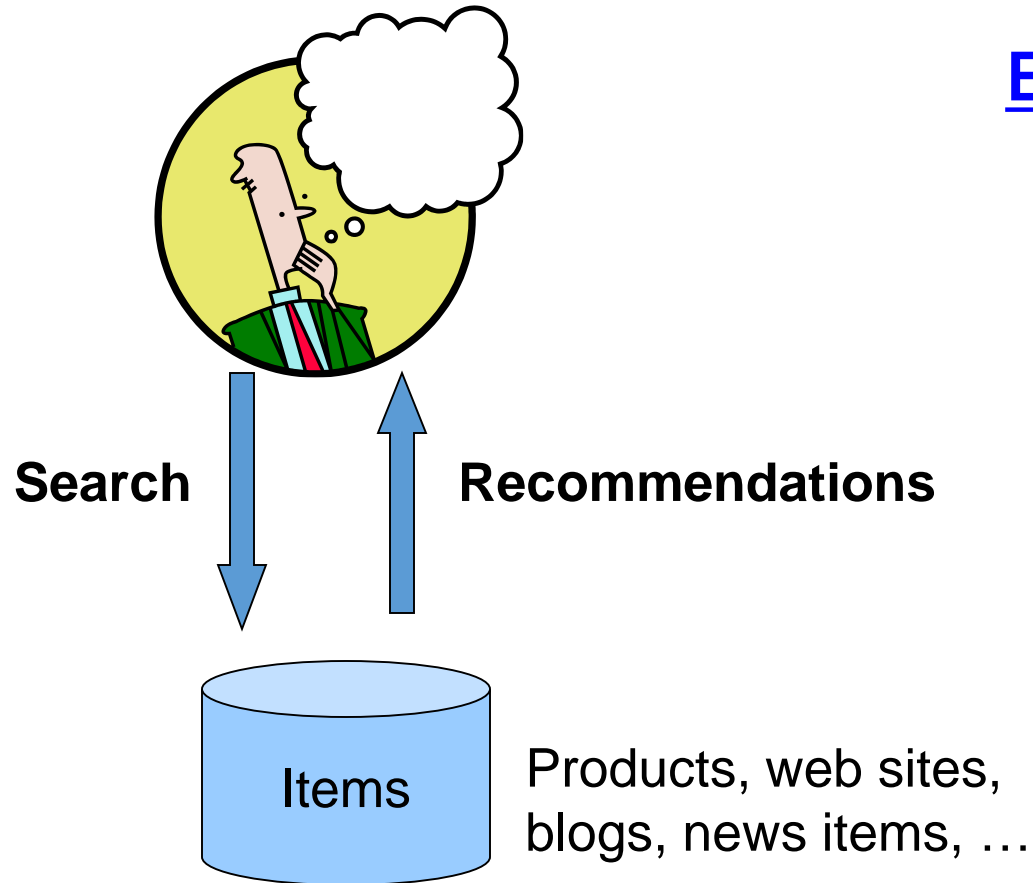
## ■ Scholar X

- Buys Metallica CD
- Buys Megadeth CD

## ■ Scholar Y

- Does search on Metallica
- Recommender system suggests Megadeth from data collected about customer X

# Recommendations



## Examples:

amazon.com.




**m o v i e l e n s**  
helping you find the *right* movies



# From Scarcity to Abundance

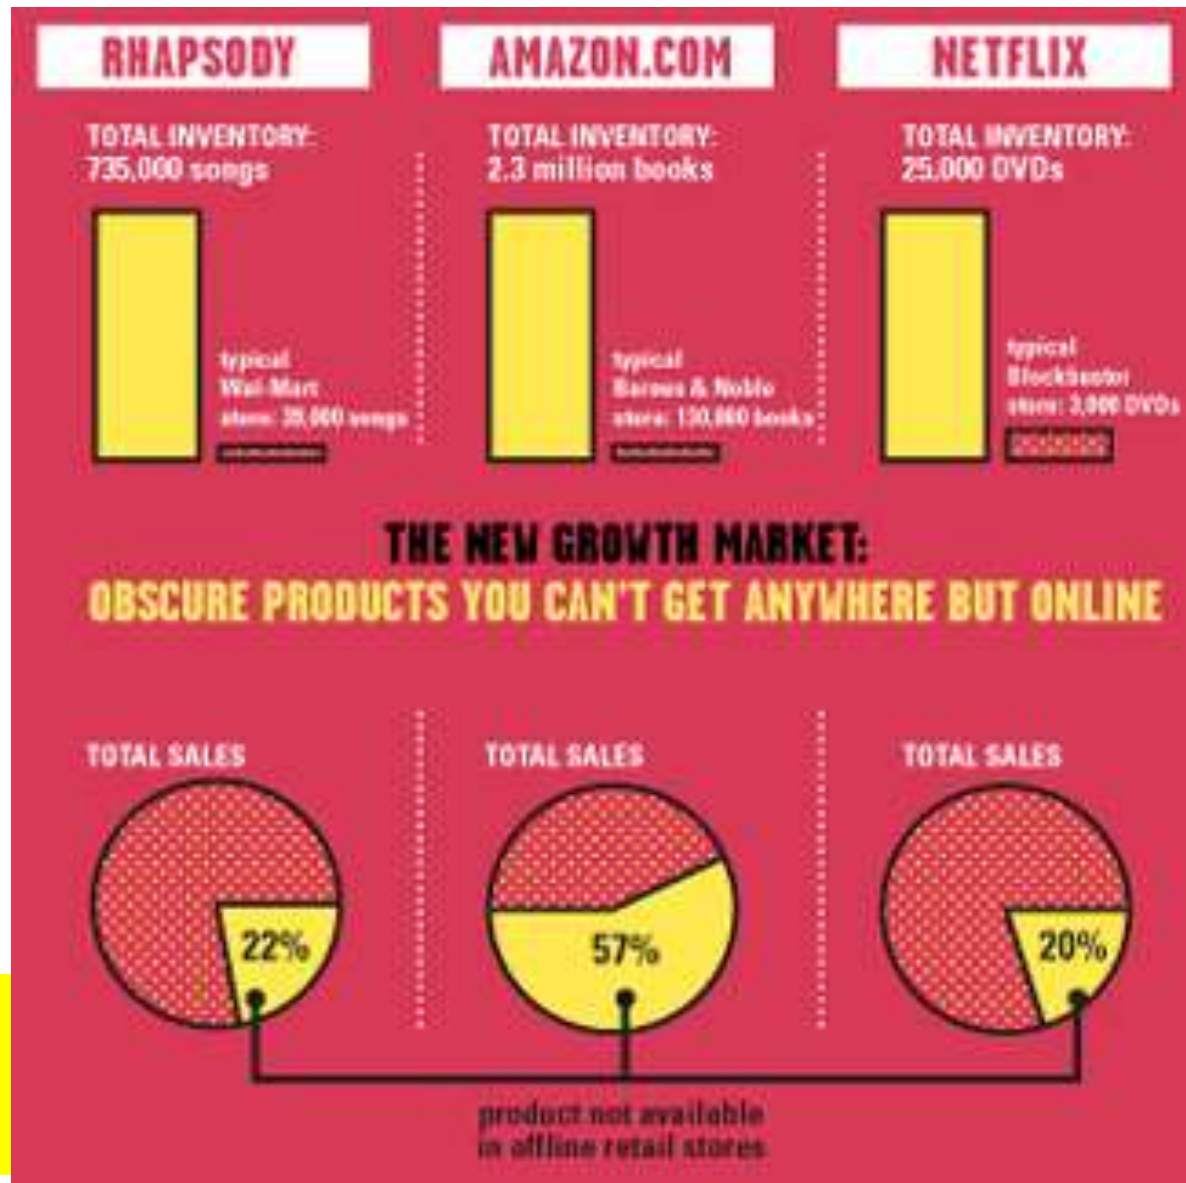
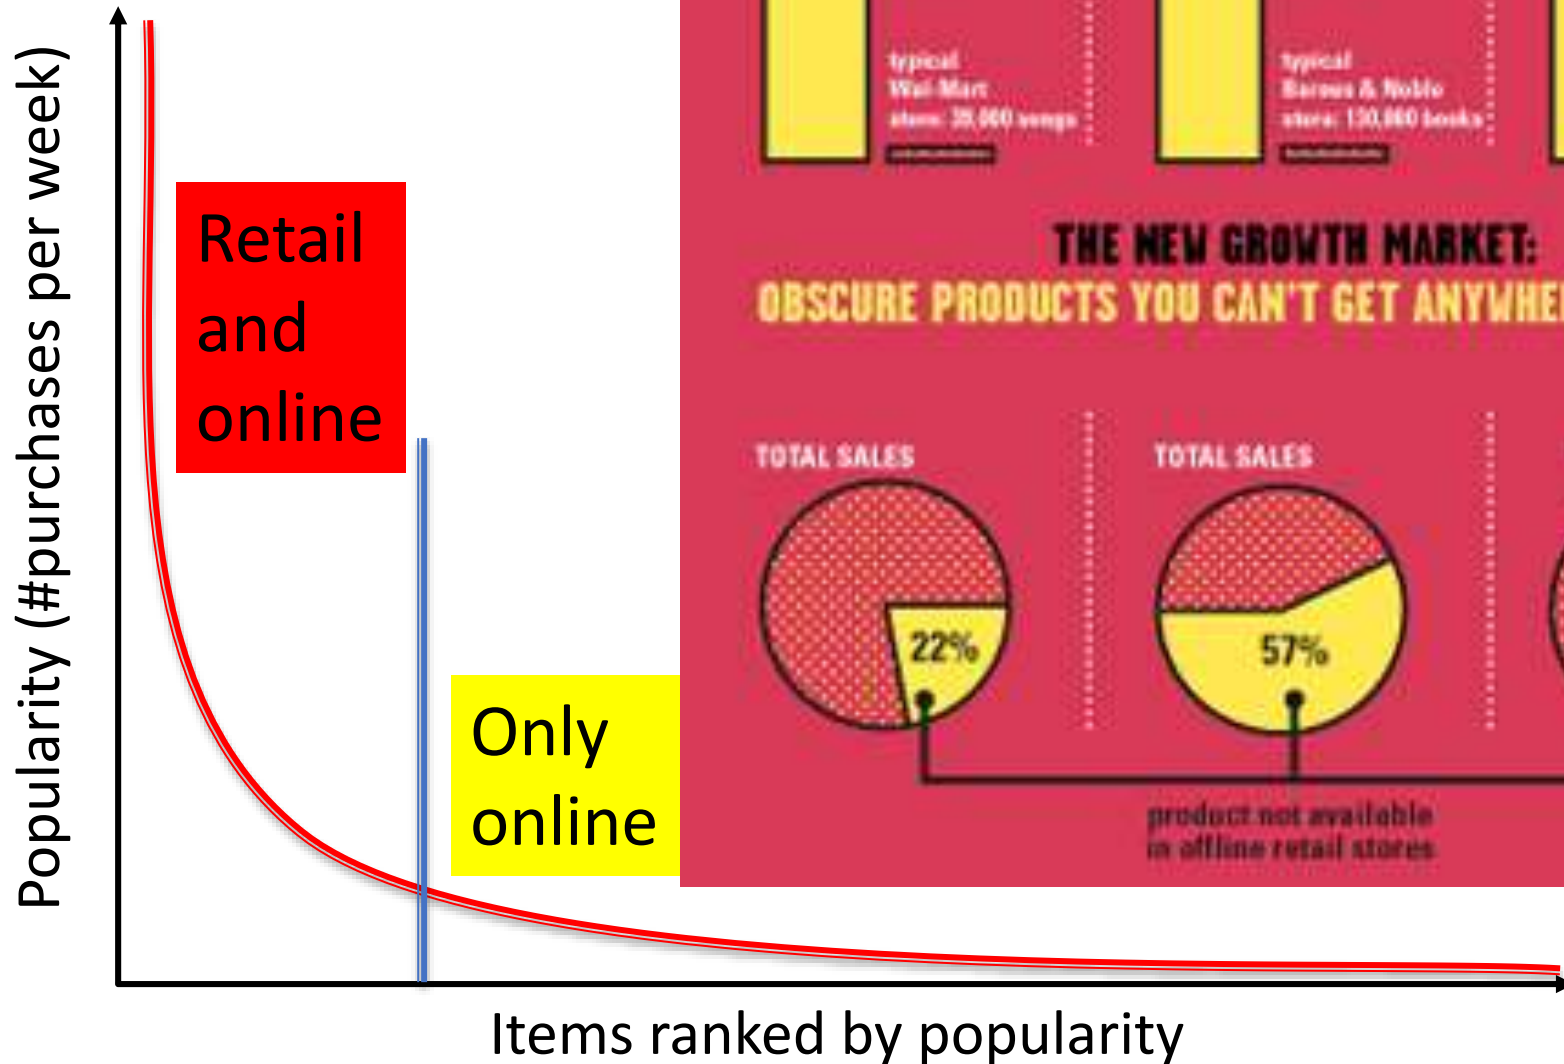
- **Shelf space is a scarce commodity for traditional retailers**
  - Also: TV networks, movie theaters,...
- **Web enables near-zero-cost dissemination of information about products**
  - From scarcity to abundance
- More choice necessitates **better filters**
  - Recommendation engines
  - How **Into Thin Air** made **Touching the Void** a bestseller: <http://www.wired.com/wired/archive/12.10/tail.html>

# The Long Tail





# The Long Tail



# Formal Model

- $X$  = set of **Customers**
- $S$  = set of **Items**
- **Utility function**  $u: X \times S \rightarrow R$ 
  - $R$  = set of ratings
  - $R$  is a totally ordered set
  - e.g., **0-5** stars, real number in **[0,1]**

# Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

# Key Problems

- (1) **Gathering “known” ratings for matrix**
  - How to collect the data in the utility matrix
- (2) **Extrapolate unknown ratings from the known ones**
  - Mainly interested in high unknown ratings
    - We are not interested in knowing what you don't like but what you like
- (3) **Evaluating extrapolation methods**
  - How to measure success/performance of recommendation methods

# (1) Gathering Ratings

## ■ Explicit

- Ask people to rate items
- Doesn't work well in practice – people can't be bothered

## ■ Implicit

- Learn ratings from user actions
  - E.g., purchase implies high rating
- What about low ratings?

## (2) Extrapolating Utilities

- **Key problem:** Utility matrix  $U$  is sparse
  - Most people have not rated most items
  - **Cold start:**
    - New items have no ratings
    - New users have no history
- Three approaches to recommender systems:
  - 1) **Content-based**
  - 2) **Collaborative**
  - 3) Latent factor based

Today

# **Content-based Recommender Systems**

# Content-based Recommendations

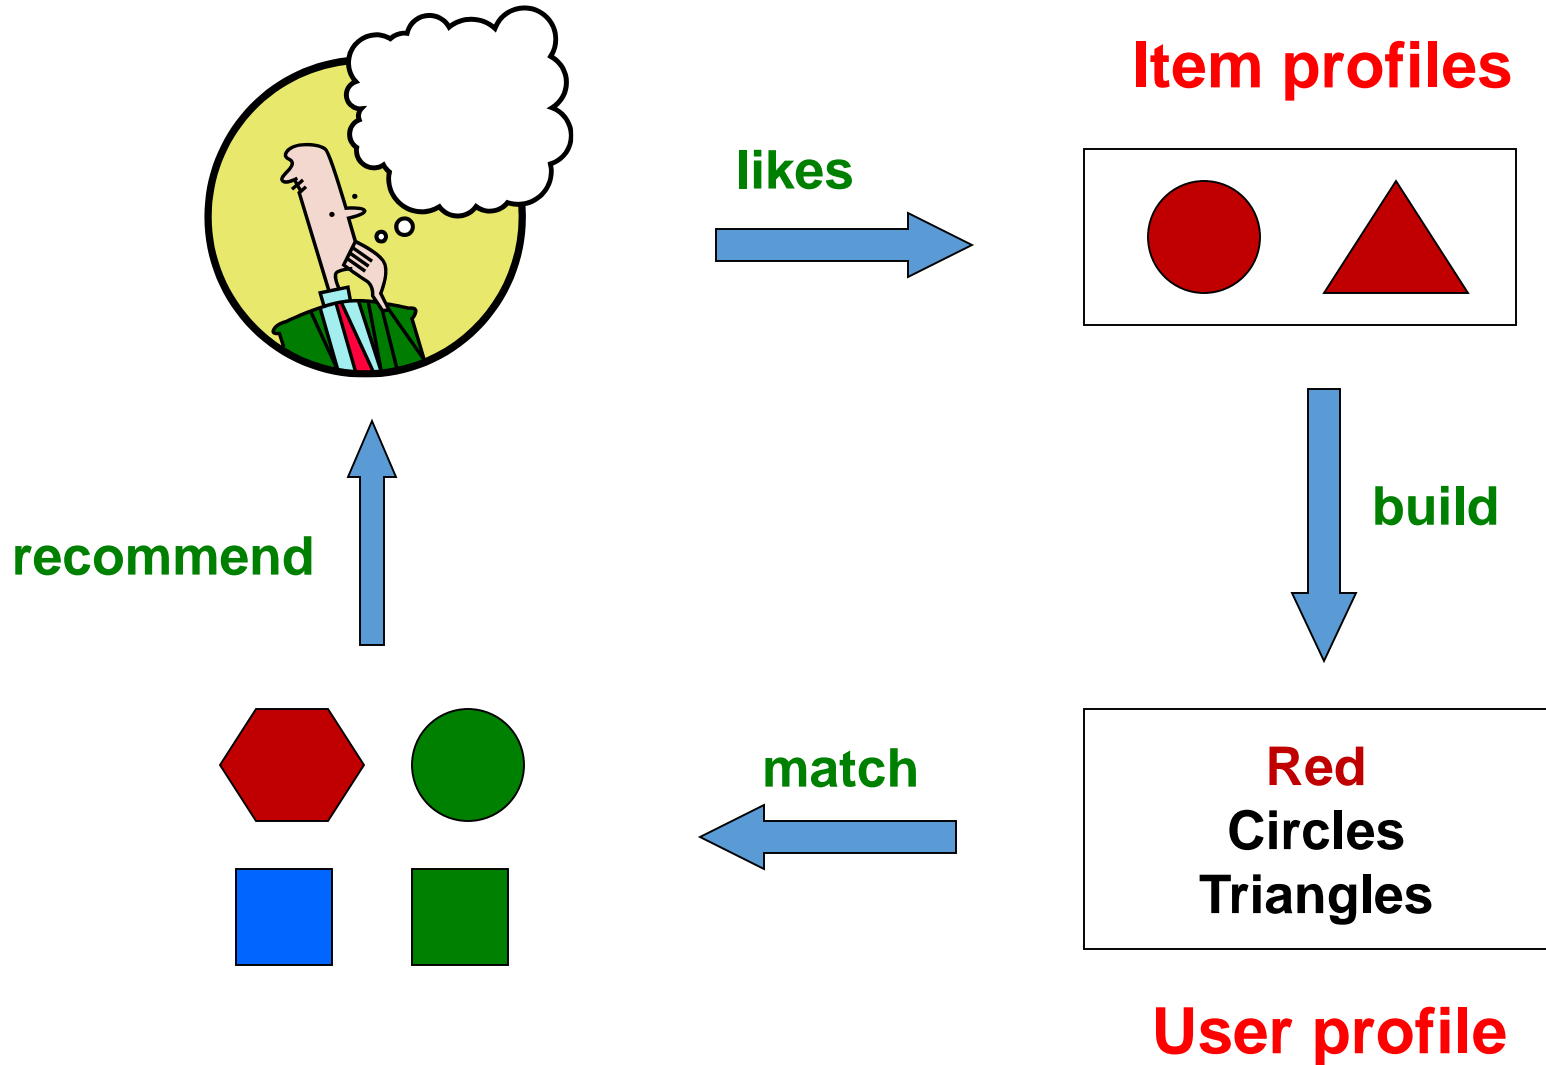
- **Main idea:** Recommend to customer  $x$  items similar to previous items rated highly by  $x$

## *Example:*

- **Movie recommendations**
  - Recommend movies with same actor(s), director, genre, ...
- **Websites, blogs, news**
  - Recommend other sites with “similar” content



# Plan of Action



# Item Profiles

- For each item, create an **item profile**
- Profile is a **set (vector) of features**
  - **Movies:** author, title, actor, director,...
  - **Text:** Set of “important” words in document
- Example Text: How to pick important features?
  - Usual heuristic from text mining is **TF-IDF**  
(**Term-frequency** \* **Inverse-Doc-Frequency**)
    - Term == Feature
    - Doc(ument) == Item

# Sidenote: **TF-IDF**

- Let  $f_{ij}$  be frequency of term  $i$  in document  $j$

- Then:

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}}$$

**Note:** we normalize TF  
to discount for “longer” docs

- Let  $n_i$  = number of docs that mention term  $i$

- And  $N$  = total number of docs

- Then:

$$IDF_i = \log \frac{N}{n_i}$$

The **TF-IDF score** is (term  $i$ , doc  $j$ ):  $w_{ij} = TF_{ij} \times IDF_i$

⇒ Use this to define a (variant of) **doc profile**:

= Set of words with highest TF-IDF scores (together with these scores)

# Example: Items are Movies

- Representing item profile – a “mixed” vector

$$i_1 = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 1 \\ \vdots \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ \vdots \\ 3.2 \\ 2.9 \\ \vdots \end{bmatrix} \begin{array}{l} \left. \begin{array}{c} \phantom{0} \\ \phantom{1} \\ \phantom{0} \\ \phantom{1} \\ \phantom{1} \\ \vdots \\ \phantom{0} \\ \phantom{0} \\ \phantom{1} \\ \phantom{0} \\ \phantom{0} \\ \vdots \end{array} \right\} \begin{array}{l} \text{Set of actors} \\ \text{(as 0/1 entries)} \end{array} \\ \left. \begin{array}{c} \phantom{0} \\ \phantom{0} \\ \phantom{1} \\ \phantom{0} \\ \phantom{0} \\ \vdots \end{array} \right\} \begin{array}{l} \text{Director(s)} \\ \text{(as 0/1 entries)} \end{array} \\ \left. \begin{array}{c} \phantom{3.2} \\ \phantom{2.9} \\ \vdots \end{array} \right\} \begin{array}{l} \text{Ratings from various} \\ \text{movie DBs (as numbers)} \end{array} \end{array}$$

# User Profiles – 1<sup>st</sup> Attempt

In general, does not work!

- Intuition: average the profiles of all items rated by a user and weight them by the ratings of this user
- User  $u$  gives items 1, 2, ...,  $n$  ratings  $r_1, r_2, \dots, r_n$
- **User profile  $x$** 
  - $x$  = weighted average of rated item profiles

$$\mathbf{x} = (r_1 * i_1 + r_2 * i_2 + \dots + r_n * i_n) / n$$

Weights = ratings

Profiles of items 1, 2, ..,  $n$  (those rated by user  $u$ ) - vectors

# Example: Star-Based Ratings /1

$$i_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, i_2 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, i_3 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}, i_4 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}, i_5 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

← actor A present  
← actor B present

- Items are movies, only features are “actors”
  - Item profile has 2 components (for actor A and actor B)
- User ratings are 1 to 5 stars (per movie)
- User watched 5 movies
  - Actor A – movies got 3 and 5 stars (movies 1 & 2)
  - Actor B – movies got 1, 2 and 4 stars (movies 3, 4, 5)
- Ratings are  $r_1=3, r_2=5, r_3=1, r_4=2, r_5=4$
- Item profiles are as above

# Example: Star-Based Ratings /2

- The user profile becomes:

$$(r_1 i_1 + r_2 i_2 + r_3 i_3 + r_4 i_4 + r_5 i_5) / 5 =$$

$$\left( 3 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 5 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 1 \begin{bmatrix} 0 \\ 1 \end{bmatrix} + 2 \begin{bmatrix} 0 \\ 1 \end{bmatrix} + 4 \begin{bmatrix} 0 \\ 1 \end{bmatrix} \right) / 5 = \begin{bmatrix} 8/5 \\ 7/5 \end{bmatrix}$$

- Problem 1: user likes actor A more than actor B, but this shows only weakly in his profile!
- Problem 2: with more ratings, each component becomes smaller (as **n** gets larger)
  - Because components with value 0 disturb the average, but should be treated as “don’t care about corresp. rating”

# For 1: Normalizing Ratings

- Solution for 1: Normalize ratings by subtracting user's mean rating (which is  $3 = (3+5+1+2+4)/5$ )
  - Normalized ratings for actor A movies  $\Rightarrow 0, +2$
  - Normalized ratings for actor A movies  $\Rightarrow -2, -1, +1$
- Then the user profile is:
  - With  $r_1=0, r_2=+2, r_3=-2, r_4=-1, r_5=+1$

$$\left(0 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 2 \begin{bmatrix} 1 \\ 0 \end{bmatrix} - 2 \begin{bmatrix} 0 \\ 1 \end{bmatrix} - 1 \begin{bmatrix} 0 \\ 1 \end{bmatrix} + 1 \begin{bmatrix} 0 \\ 1 \end{bmatrix}\right) / 5 = \begin{bmatrix} 2/5 \\ -2/5 \end{bmatrix}$$

Now better: clear distinction  
for actor A and actor B



# For 2: Per-component Weights

Only ratings for these 2 items should be counted for actor A

$$\left(0 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 2 \begin{bmatrix} 1 \\ 0 \end{bmatrix} - 2 \begin{bmatrix} 0 \\ 1 \end{bmatrix} - 1 \begin{bmatrix} 0 \\ 1 \end{bmatrix} + 1 \begin{bmatrix} 0 \\ 1 \end{bmatrix}\right) / 5 = \begin{bmatrix} 2/5 \\ -2/5 \end{bmatrix}$$

Only the ratings for these 3 items  
should be counted for actor B

- Essence of problem 2: a 0 in an item's component (=attribute) k should mean “don't care”, but now mean “one more neutral rating for attribute k”
- => Use “individual”  $n$  for each vector component
  - For actor A:  $n_A = 2$ , for actor B:  $n_B = 3$

# For 2: Per-component Weights

- => Use “individual”  $n$  for each vector component
  - For actor A:  $n_A = 2$ , for actor B:  $n_B = 3$
  - Recall: Normalized ratings are  $r_1=0$ ,  $r_2=+2$ ,  $r_3=-2$ ,  $r_4=-1$ ,  $r_5=+1$
- Then the user profile becomes:


$$\begin{bmatrix} r_1 / n_a \\ 0 \end{bmatrix} + \begin{bmatrix} r_2 / n_a \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ r_3 / n_b \end{bmatrix} + \begin{bmatrix} 0 \\ r_4 / n_b \end{bmatrix} + \begin{bmatrix} 0 \\ r_5 / n_b \end{bmatrix} =$$

$$\begin{bmatrix} 0/2 \\ 0 \end{bmatrix} + \begin{bmatrix} 2/2 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ -2/3 \end{bmatrix} + \begin{bmatrix} 0 \\ -1/3 \end{bmatrix} + \begin{bmatrix} 0 \\ 1/3 \end{bmatrix} = \begin{bmatrix} 1 \\ -2/3 \end{bmatrix}$$

# Matching User and Item Profiles

- To compute similarity of user profile and item profile, use a **prediction heuristic**:
  - Given **user profile  $x$**  and **item profile  $i$** , estimate

$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{||x|| \cdot ||i||}$$




# Summary: Content-Based R.

- We construct a vector  $i$  **for each item** (“**item profile**”) and a vector  $x$  (of size  $s$ ) **for each user**
  - Item profile  $i$ : “natural” attributes of an item
  - User vector  $x$ : combination of item profiles rated by this user

- **Prediction heuristic:**

- Given a user vector  $x$  and item vector  $i$ , estimate similarity

$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{||x|| \cdot ||i||}$$



- For a user with vector  $x$ , recommend by various criteria:
    - E.g. all items with  $u(x, i) > \text{threshold}$
    - Rank items by  $u(x, i)$ , recommend top  $k$  (e.g.  $k=5$ )

# Pros: Content-based Approach

- **+: No need for data on other users**
  - No cold-start or sparsity problems
- **+: Able to recommend to users with unique tastes**
- **+: Able to recommend new & unpopular items**
  - No first-rater problem
- **+: Able to provide explanations**
  - Can provide explanations of recommended items by listing content-features that caused an item to be recommended

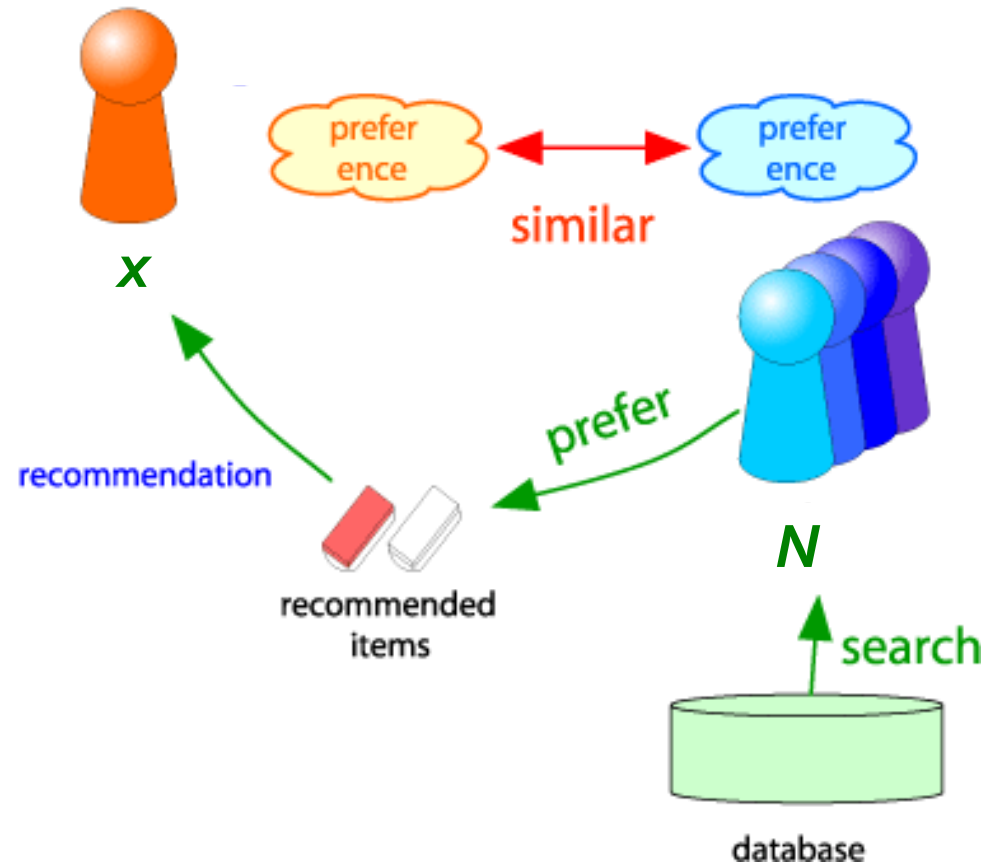
# Cons: Content-based Approach

- —: **Finding the appropriate features is hard**
  - E.g., images, movies, music
- —: **Recommendations for new users**
  - **How to build a user profile?**
- —: **Overspecialization**
  - Never recommends items outside user's content profile
  - People might have multiple interests
  - **Unable to exploit quality judgments of other users**

# Collaborative Filtering

# User-User Collaborative Filtering

- Consider user  $x$
- Find set  $N$  of other users whose ratings are “**similar**” to  $x$ ’s ratings
- Estimate  $x$ ’s ratings based on ratings of users in  $N$





# Finding “Similar” Users

- Let  $r_x$  be the vector of user  $x$ 's ratings

$$\begin{aligned} r_x &= [*, \_, \_, *, **] \\ r_y &= [*, \_, **, **, \_] \end{aligned}$$

- **Cosine similarity measure**

- $\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{||r_x|| \cdot ||r_y||}$

*$r_x, r_y$  as points:*

$$\begin{aligned} r_x &= \{1, 0, 0, 1, 3\} \\ r_y &= \{1, 0, 2, 2, 0\} \end{aligned}$$

- **Problem:** Treats missing ratings as “negative”


# Jaccard Measures

- The **Jaccard similarity** of two **sets** is the size of their intersection divided by the size of their union:

$$\text{sim}(\mathbf{C}_1, \mathbf{C}_2) = |\mathbf{C}_1 \cap \mathbf{C}_2| / |\mathbf{C}_1 \cup \mathbf{C}_2|$$

- **Jaccard distance:**  $d(\mathbf{C}_1, \mathbf{C}_2) = 1 - |\mathbf{C}_1 \cap \mathbf{C}_2| / |\mathbf{C}_1 \cup \mathbf{C}_2|$

- For measuring similarity of users, we consider only sets of items for which users voted
- Problem? Values of ratings are ignored!



3 in intersection  
8 in union  
Jaccard similarity =  $3/8$   
Jaccard distance =  $5/8$

$$\begin{aligned} r_x &= [* , \_ , \_ , * , ***] \\ r_y &= [* , \_ , ** , ** , \_] \end{aligned}$$

$$\begin{aligned} &r_x, r_y \text{ as sets:} \\ r_x &= \{1, 4, 5\} \\ r_y &= \{1, 3, 4\} \end{aligned}$$

# A Good Similarity Metric

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Intuitively we want:  **$\text{sim}(A, B) > \text{sim}(A, C)$**

- Jaccard similarity:  $1/5 < 2/4 \Rightarrow$  bad

- Cosine similarity:  $0.386 > 0.322 \Rightarrow$  not good

- Solution: subtract the (row) mean**

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

Notice: cos similarity  
is = **correlation**  
when data is  
centered at 0!

**sim A,B vs. A,C:**  
 **$0.092 > -0.559$**

$$\Rightarrow \text{sim}(x, y) = \frac{\sum_i r_{xi} \cdot r_{yi}}{\sqrt{\sum_i r_{xi}^2} \cdot \sqrt{\sum_i r_{yi}^2}}$$

# From Cosine to Pearson

- Let  $\mathbf{r}_x$  be the vector of user  $\mathbf{x}$ 's ratings

- **Cosine similarity measure**

- $\text{sim}(\mathbf{x}, \mathbf{y}) = \cos(\mathbf{r}_x, \mathbf{r}_y) = \frac{\mathbf{r}_x \cdot \mathbf{r}_y}{\|\mathbf{r}_x\| \cdot \|\mathbf{r}_y\|}$

- **Pearson correlation coefficient**

- $S_{xy}$  = items rated by both users  $\mathbf{x}$  and  $\mathbf{y}$

$\bar{r}_x, \bar{r}_y \dots$  avg.  
rating of  $\mathbf{x}, \mathbf{y}$

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\sum_{s \in S_{xy}} (\mathbf{r}_{xs} - \bar{r}_x)(\mathbf{r}_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (\mathbf{r}_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (\mathbf{r}_{ys} - \bar{r}_y)^2}}$$

$$\mathbf{r}_x = [* , \_ , \_ , * , ***]$$
$$\mathbf{r}_y = [* , \_ , ** , ** , \_]$$

# Rating Predictions

From similarity metric to recommendations:

- Let  $\mathbf{r}_x$  be the vector of user  $x$ 's ratings
- Let  $N$  be the set of  $k$  users most similar to  $x$  (according to  $\text{sim}(x, y)$ ) who have rated item  $i$
- Prediction for item  $i$  of user  $x$ :

- $$r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$$

Users rating of item  $i$  := avg. rating of  $k$  most similar users

- $$r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$$

Better: weights a recommendation of other user  $y$  by similarity  $s_{xy}$  to this "neighbor"  $y$

Shorthand:

$$s_{xy} = \text{sim}(x, y)$$

# Item-Item Collaborative Filtering

## ■ Another view: Item-item

- For item  $i$ , find other **similar items** rated by user  $x$
- Estimate rating for item  $i$  based on ratings for similar items (with ratings by user  $x$ )
  - Note: We can use same similarity metrics and prediction functions as in user-user model

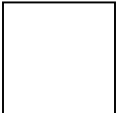
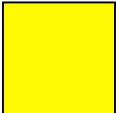
$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

Shorthand:


$$s_{ij} = \mathbf{sim}(i, j)$$

$r_{xj}$  rating of user  $x$  on item  $j$   
 $N(i;x)$  set of items rated by  $x$  similar to  $i$

# Example: Item-Item CF ( $|N|=2$ )

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3			5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	
		 - unknown rating		 - rating between 1 to 5									

# Item-Item CF ( $|N|=2$ )



- estimate rating of movie **1** by user **5**

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



# Item-Item CF ( $|N|=2$ )

## Neighbor selection:

Identify **movies similar to movie 1**, rated by user 5

We use Pearson correlation as similarity:

1) Subtract mean rating  $m_i$  from each movie  $i$

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1:  $[-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]$

2) Compute cosine similarities between rows

users

	1	2	3	4	5	6	7	8	9	10	11	12	sim(1,m)
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	
<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
<u>6</u>	1		3		3			2			4		<u>0.59</u>

# Item-Item CF ( $|N|=2$ )

Compute similarity weights:

$$s_{1,3} = \mathbf{0.41}, s_{1,6} = \mathbf{0.59}$$

users

	1	2	3	4	5	6	7	8	9	10	11	12	sim(1,m)
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	-
<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
<u>6</u>	1		3		3			2			4		<u>0.59</u>

movies

# Item-Item CF ( $|N|=2$ )

$$r_{ix} = \frac{\sum_{j \in N(i,x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$

Predict by taking weighted average:

$$r_{1,5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	sim(1,m)
movies	1	1		3		2.6	5			5		4		1.00
	2			5	4			4			2	1	3	
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		<u>0.59</u>

# CF: Common Practice

Before (simpler):

- Define similarity  $s_{ij}$  of items  $i$  and  $j$
- Select  $k$  nearest neighbors  $N(i; x)$ 
  - Items most similar to  $i$ , that were rated by  $x$
- Estimate rating  $r_{xi}$  as the weighted average:

$$r_{xi} = \frac{\sum_{j \in N(i; x)} s_{ij} r_{xj}}{\sum_{j \in N(i; x)} s_{ij}}$$

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i; x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i; x)} s_{ij}}$$

baseline estimate for  $(x, i)$

baseline estimate for  $(x, j)$

baseline estimate for  $(x, i)$ :

$$b_{xi} = \mu + b_x + b_i$$

- $\mu$  = overall mean item rating
- $b_x$  = rating deviation of user  $x$   
= (avg. rating of user  $x$ ) -  $\mu$
- $b_i$  = rating deviation of item  $i$

# Item-Item vs. User-User

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

- In practice, it has been observed that item-item often works better than user-user
- **Why?** Items are simpler, users have multiple tastes

# Pros/Cons of Collaborative Filtering

- **+ Works for any kind of item**
  - No feature selection needed
- **- Cold Start:**
  - Need enough users in the system to find a match
- **- Sparsity:**
  - The user/ratings matrix is sparse
  - Hard to find users that have rated the same items
- **- First rater:**
  - Cannot recommend an item that has not been previously rated (e.g. new items, esoteric items)
- **- Popularity bias:**
  - Cannot recommend items to someone with unique taste
  - Tends to recommend popular items

# Hybrid Methods

- Implement two or more different recommenders and combine predictions
  - Perhaps using a linear model
- Add content-based methods to collaborative filtering
  - Item profiles for new item problem
  - Demographics to deal with new user problem

# **Remarks & Practical Tips**



# Evaluation

**movies**

**users**

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

# Evaluation

**movies**

**users**

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

**Test Data Set**

# Evaluating Predictions

- **How to compare predictions with known ratings?**

- **Root-mean-square error (RMSE)**, details: [link](#)

- $$\sqrt{\frac{1}{N} \sum_{xi} (r_{xi} - r_{xi}^*)^2}$$

- where  $r_{xi}$  is predicted,  $r_{xi}^*$  is the true rating of  $x$  on  $i$ , and  $N$  is the number of ratings (= number of used  $(x,i)$  combinations)

- **Precision at top 10**: % of those in top 10

- Another approach: **0/1 model**

- **Coverage**:

- Number of items/users for which system can make predictions

- **Precision**:

- Accuracy of predictions

- **Receiver operating characteristic (ROC)**

- Tradeoff curve between false positives and false negatives

# Collaborative Filtering: Complexity

- Expensive step is finding  $k$  most similar customers:  $O(|X|)$
- Too expensive to do at runtime
  - Could pre-compute
- Naïve pre-computation takes time  $O(k \cdot |X|)$ 
  - $X$  ... set of customers
- We already know how to do this!
  - Near-neighbor search in high dimensions (**LSH**)
  - Clustering
  - Dimensionality reduction

# Tip: Add Data

- **Leverage all the data**

- Don't try to reduce data size in an effort to make fancy algorithms work
- Simple methods on large data do best

- **Add more data**

- e.g., add IMDB data on genres

- **More data beats better algorithms**

<http://anand.typepad.com/datawocky/2008/03/more-data-usual.html>

**Thank you.**

Questions?