

# Mining Massive Datasets

## Lecture 8

Artur Andrzejak

<http://pvs.ifi.uni-heidelberg.de>



RUPRECHT-KARLS-  
UNIVERSITÄT  
HEIDELBERG

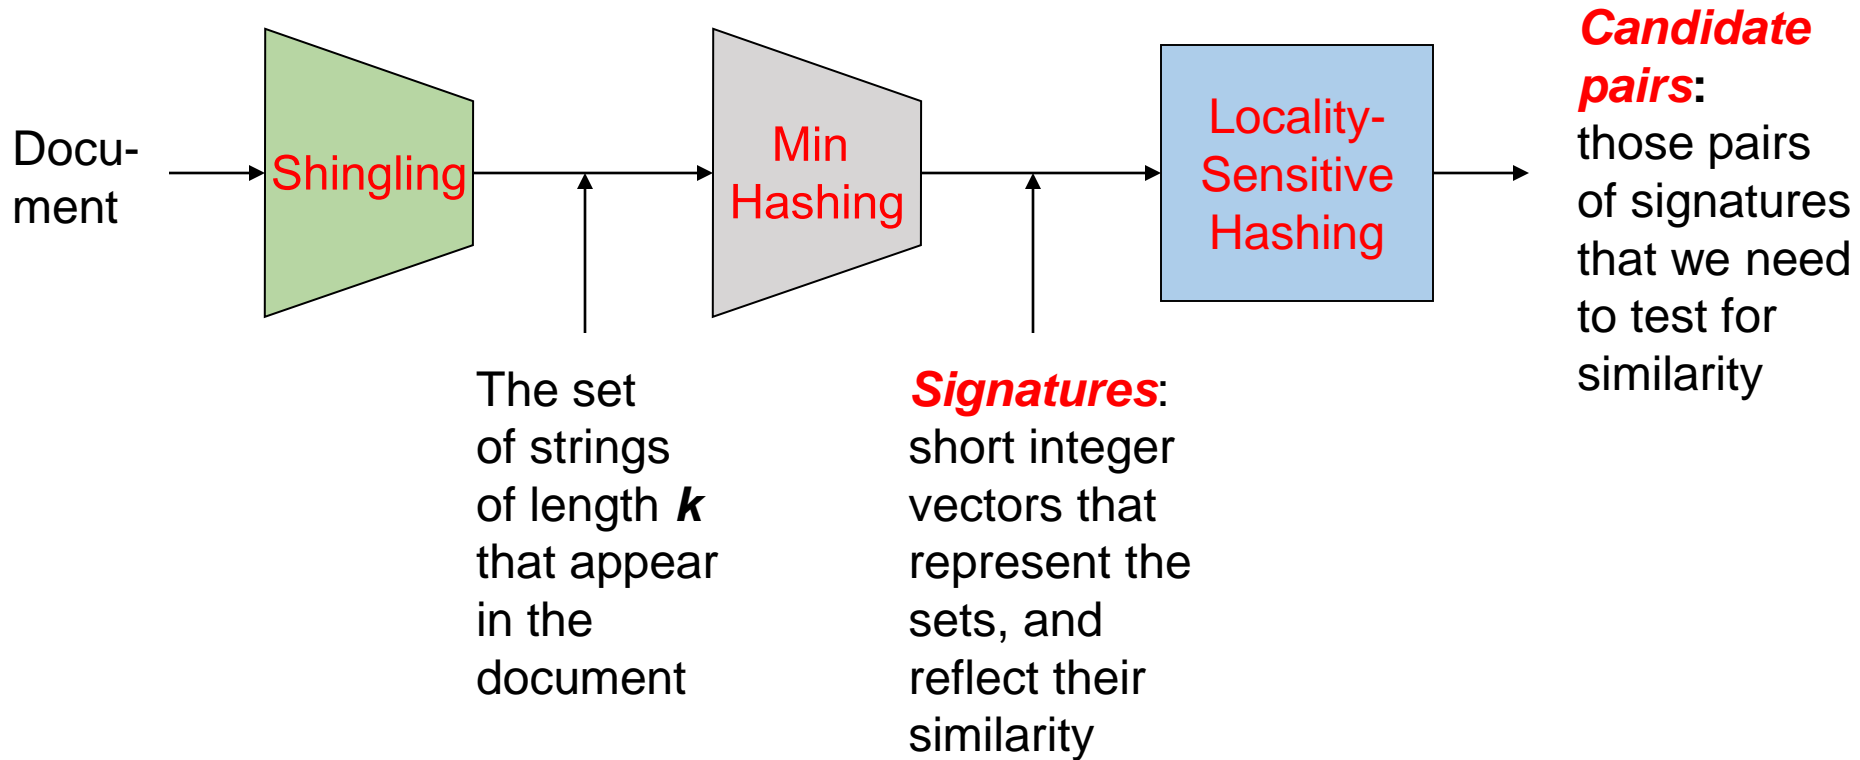


# Note on Slides

A substantial part of these slides come (either verbatim or in a modified form) from the book *Mining of Massive Datasets* by Jure Leskovec, Anand Rajaraman, Jeff Ullman (Stanford University).

For more information, see the website accompanying the book: <http://www.mmds.org>.

# Recall



# MinHashing

Finishing ...


# MinHash: Repetition

- Def.: Let  $h$  be a hash function that maps the members of  $S$  to distinct integers, then for any set  $S$  define  $\text{MinHash}_h(S) = h_{\min}(S)$  to be the minimum value of  $h(x)$

- **Example:**

- Assume  $S = \{2, 3, 6\}$  and
- $h(2) = 4$ ,  $h(3) = 5$ ,  $h(6) = 2$

- $h_{\min}(S) = 2$



# Min-Hashing: Other Interpretation

- We represent each set as a Boolean vector **C** (here:  $S = \{2, 3, 6\}$ )
- Assume that a hash function **h** is given by a (random) permutation  $\pi$  of the rows of the Boolean vector
  - **h** fulfills: "... maps the members of S to distinct integers"

Row perm.		
1	0	$\pi(1) = 3$
2	1	$\pi(2) = 4$
3	1	$\pi(3) = 5$
4	0	$\pi(4) = 6$
5	0	$\pi(5) = 1$
6	1	$\pi(6) = 2$

- Then **MinHash<sub>h</sub>**(S) is the index of the **first row** of the permuted column **C** with value **1**
- $\Rightarrow$  Again,  $h_{\pi, \min}(S) = 2$

1	
<u>2</u>	1
3	
<u>4</u>	1
<u>5</u>	1
6	

# Implementation /1

- Permuting rows even once is very expensive
- Approximate permutation by a hash function  $h_i$ 
  - $h_i(x) = [((a \cdot x + b) \bmod p) \bmod N] + 1$
  - $a, b$ : random integers;  
 $p$ : a prime ( $p > N$ );  $N$ : #rows in the matrix
  - $h_i$  is possibly not injective, but errors are rare  $\Rightarrow$  OK
- Pick about **K = 100** such hash functions  $h_i$

1	a	$h_i(1) = 3$
2	b	$h_i(2) = 4$
3	c	$h_i(3) = 1$
4	d	$h_i(4) = 2$


c
d
a
b

# Implementation /2

- Pick about  $K = 100$  such hash functions  $h_i$

Intuition: for fixed  $C$  and  $h_i$ , find the smallest value  $h_i(r)$  over all rows  $r$  with  $C(r) = 1$

## One-pass implementation:

- For each column  $C$  and hash-function  $h_i$  prepare a “slot” (variable)  $sig(C)[i]$  for the min-hash value
- Initialize all  $sig(C)[i] = \infty$
- **Scan rows looking for 1s**
  - If row  $q$  has 1 in column  $C$ , then for each  $h_i$  ( $i=1..100$ ):
    - If  $h_i(q) < sig(C)[i]$ , then  $sig(C)[i] \leftarrow h_i(q)$



# Implementation /3

For fixed  $\mathbf{C}$  and  $h_i$ :

Find the smallest value  $h_i(r)$  over all rows  $r$  with  $\mathbf{C}(r) = 1$


## Scan rows looking for 1s

- If row  $q$  has 1 in column  $\mathbf{C}$ , then for each  $h_i$ :
  - If  $h_i(q) < \text{sig}(\mathbf{C})[i]$ , then  $\text{sig}(\mathbf{C})[i] \leftarrow h_i(q)$

Example: fixed  $\mathbf{C}$  and  $h_i$

$S = \{2, 3\}$

4      1




	$\mathbf{C}$	
1	0	$h_i(1) = 3$
2	1	$h_i(2) = 4$
3	1	$h_i(3) = 1$
4	0	$h_i(4) = 2$

$\text{sig}(\mathbf{C})[i]$

$\text{sig}(\mathbf{C})[i]$
$\infty$
4
1
1

# Locality Sensitive Hashing



Step 3: **Locality-Sensitive Hashing:**  
Focus on pairs of signatures likely to be from similar documents

# LSH: First Cut

2	1	4	1
1	2	1	2
2	1	2	1

- **Goal:** Find documents with Jaccard similarity at least  $s$  (for some similarity threshold, e.g.,  $s=0.8$ )
- **LSH – General idea:** Use a function  $f(x,y)$  that tells whether  $x$  and  $y$  is a *candidate pair*: a pair of elements whose similarity must be evaluated
- **For Min-Hash matrices:**
  - Hash columns of *signature matrix*  $M$  to many buckets
  - Each pair of documents that hashes into the same bucket is a candidate pair

# Candidates from Min-Hash

2	1	4	1
1	2	1	2
2	1	2	1

- Pick a similarity threshold  $s$  ( $0 < s < 1$ )
- Columns  $\mathbf{x}$  and  $\mathbf{y}$  of  $\mathbf{M}$  are a **candidate pair** if their signatures agree on at least fraction  $s$  of their rows:  
 $\mathbf{M}(i, \mathbf{x}) = \mathbf{M}(i, \mathbf{y})$  for at least frac.  $s$  values of  $i$ 
  - We expect documents  $\mathbf{x}$  and  $\mathbf{y}$  to have the same (Jaccard) similarity as their signatures


# LSH for Min-Hash

2	1	4	1
1	2	1	2
2	1	2	1

- **Big idea:** Hash columns of signature matrix  $M$  several times
- Arrange that (only) **similar columns** are likely to **hash to the same bucket**, with high probability
- **Candidate pairs are those that hash to the same bucket**

# Partition $M$ into $b$ Bands


2	1	4	1
1	2	1	2
2	1	2	1



# Partition $M$ into Bands

- Divide matrix  $M$  into  $b$  bands of  $r$  rows
- For each band, hash its portion of each column to a hash table with  $k$  buckets
  - Make  $k$  as large as possible
- **Candidate** column pairs are those that hash to the same bucket for  $\geq 1$  band
- Tune  $b$  and  $r$  to catch most similar pairs, but few non-similar pairs

# Hashing Bands





# Simplifying Assumption

- There are **enough buckets** that columns are unlikely to hash to the same bucket unless they are **identical** in a particular band
- Hereafter, we assume that “**same bucket**” means “**identical in that band**”
- Assumption needed only to simplify analysis, not for correctness of algorithm

# Example of Bands

2	1	4	1
1	2	1	2
2	1	2	1

## Assume the following case:

- Suppose 100,000 columns of  $\mathbf{M}$  (100k docs)
- Signatures of 100 integers (rows)
- Therefore, signatures take 40Mb
- Choose  $\mathbf{b} = 20$  bands of  $\mathbf{r} = 5$  integers/band
- **Goal:** Find pairs of documents that are at least  $\mathbf{s} = 0.8$  similar

# $C_1, C_2$ are 80% Similar

2	1	4	1
1	2	1	2
2	1	2	1

- Find pairs of  $\geq s=0.8$  similarity, set  $b=20$ ,  $r=5$
- **Assume:**  $\text{sim}(C_1, C_2) = 0.8$ 
  - Since  $\text{sim}(C_1, C_2) \geq s$ , we want  $C_1, C_2$  to be a **candidate pair**: We want them to hash to at **least 1 common bucket** (at least one band is identical)
- **Probability  $C_1, C_2$  identical in one particular band:**  $(0.8)^5 = 0.328$
- Probability  $C_1, C_2$  are **not** similar in all of the 20 bands:  $(1-0.328)^{20} = 0.00035$ 
  - i.e., about 1/3000th of the 80%-similar column pairs are **false negatives** (we miss them)
  - **We would find 99.965% pairs of truly similar documents**

# $C_1, C_2$ are 30% Similar

2	1	4	1
1	2	1	2
2	1	2	1


- Find pairs of  $\geq s=0.8$  similarity, set  $b=20$ ,  $r=5$
- **Assume:**  $\text{sim}(C_1, C_2) = 0.3$ 
  - Since  $\text{sim}(C_1, C_2) < s$  we want  $C_1, C_2$  to hash to **NO common buckets** (all bands should be different)
- **Probability  $C_1, C_2$  identical in one particular band:**  $(0.3)^5 = 0.00243$
- Probability  $C_1, C_2$  identical in at least 1 of 20 bands:  $1 - (1 - 0.00243)^{20} = 0.0474$ 
  - In other words, approximately 4.74% pairs of docs with similarity 0.3% end up becoming **candidate pairs**
    - They are **false positives** since we will have to examine them (they are candidate pairs) but then it will turn out their similarity is below threshold  $s$

# LSH Involves a Tradeoff


2	1	4	1
1	2	1	2
2	1	2	1

- **Pick:**
  - The number of Min-Hashes (rows of  $\mathbf{M}$ )
  - The number of bands  $\mathbf{b}$ , and
  - The number of rows  $\mathbf{r}$  per bandto balance false positives/negatives
- **Example:** If we had only 15 bands of 5 rows, the number of false positives would go down, but the number of false negatives would go up

# Analysis of LSH – What We Want



# What 1 Band of 1 Row Gives You




# $b$ bands, $r$ rows/band

- Columns  $C_1$  and  $C_2$  have similarity  $t$
- Pick any band ( $r$  rows)
  - Prob. that all rows in band equal =  $t^r$
  - Prob. that some row in band unequal =  $1 - t^r$
- Prob. that no band identical =  $(1 - t^r)^b$
- Prob. that at least 1 band identical =  
 $1 - (1 - t^r)^b$



# What $b$ Bands of $r$ Rows Gives You



Example:  $b = 20; r = 5$


- **Similarity threshold  $s$**
- **Prob. that at least 1 band is identical:**

$s$	$1-(1-s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

# Picking $r$ and $b$ : The S-curve

- Picking  $r$  and  $b$  to get the best S-curve

- 50 hash-functions ( $r=5$ ,  $b=10$ )



# LSH Summary

- Tune  $M, b, r$  to get almost all pairs with similar signatures, but eliminate most pairs that do not have similar signatures
- Check in main memory that **candidate pairs** really do have **similar signatures**
- **Optional:** In another pass through data, check that the remaining candidate pairs really represent similar documents

# Summary: 3 Steps

- **Shingling:** Convert documents to sets
  - We used hashing to assign each shingle an ID
- **Min-Hashing:** Convert large sets to short signatures, while preserving similarity
  - We used **similarity preserving hashing** to generate signatures with property  $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
  - We used hashing to get around generating random permutations
- **Locality-Sensitive Hashing:** Focus on pairs of signatures likely to be from similar documents
  - We used hashing to find **candidate pairs** of similarity  $\geq s$

# **Spark: DataFrames and Datasets**

*Repetition*

# DataFrames in Spark

- **DataFrames** (DF) are tables with named and typed data columns
  - Similar to a dataframe in R, or Pandas (Python), or tables in DBMS/SQL
  - Impose a structure and schema on data
- Example

	Time (Str)	Site (Str)	Req (Int)	Time (Str)	Site (Str)	Req (Int)	Time (Str)	Site (Str)	Req (Int)	Time (Str)	Site (Str)	Req (Int)
Row 1	ts	m	1304	ts	d	3901	ts	m	1172	ts	m	2538
Row 2	ts	d	2237	ts	d	2491	ts	m	2137	ts	d	2837
Row 3	ts	m	1600	ts	d	2288	ts	d	3176	ts	d	3400
	Partition 1			Partition 2			Partition 3			Partition 4		

# DataFrames from RDDs

```
// Read file with rows: <name, age>
```

```
filePath = „/home/immd-user/spark-2 .../examples/src/main/resources/people.txt“
```

```
parts = sc.textFile(filePath).map( lambda line: line.split(“,”) )
```

```
// Each row should become a tuple (name, age)
```

```
peopleRDD = parts.map( lambda p: (p[0], p[1].strip() )
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.appName(„Exmpl“).getOrCreate()
```

```
schema = StructType( [  
    StructField(“name“, StringType(), True),  
    StructField( „age“, StringType(), True) ] )
```

```
dfPeople = spark.createDataFrame (peopleRDD, schema)
```

```
print ( dfPeople.take(5) )
```



# SQL: More User-Friendly

Standard DF API:

# Select people older than 25

```
dfPeople.filter(dfPeople['age'] > 25).show()
```

```
# |age|name|
```

```
# | 30|Andy|
```

Same result with **SQL**:

# Register the DataFrame as a SQL temporary view

```
dfPeople.createOrReplaceTempView("people")
```

```
sqlDF = spark.sql("SELECT * FROM people where age > 25")
```

```
sqlDF.show()
```

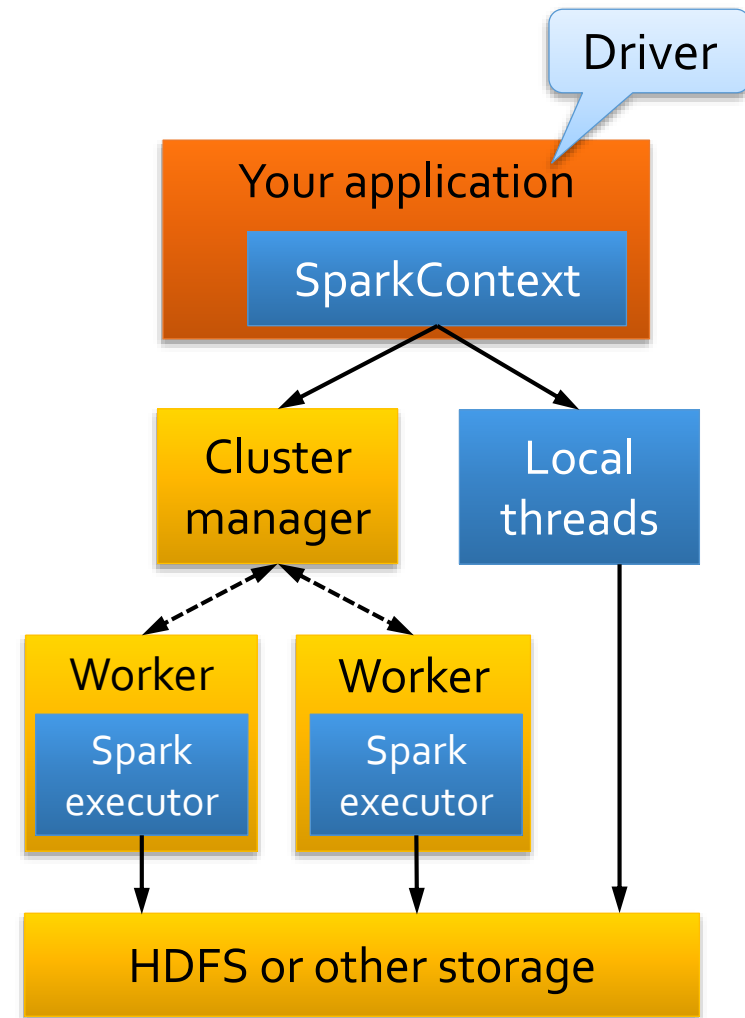
# Video

- Coursera: Big Data Integration and Processing
  - <https://www.coursera.org/learn/big-data-integration-processing/home/info>
  - Week 5: Programming in Spark
  - Video: Hands-on: Data Processing in Spark => Exploring SparkSQL and Spark DataFrames
    - Link:
      - <https://www.coursera.org/learn/big-data-integration-processing/lecture/aHc8E/exploring-sparksql-and-spark-dataframes>

# Spark: Execution Details

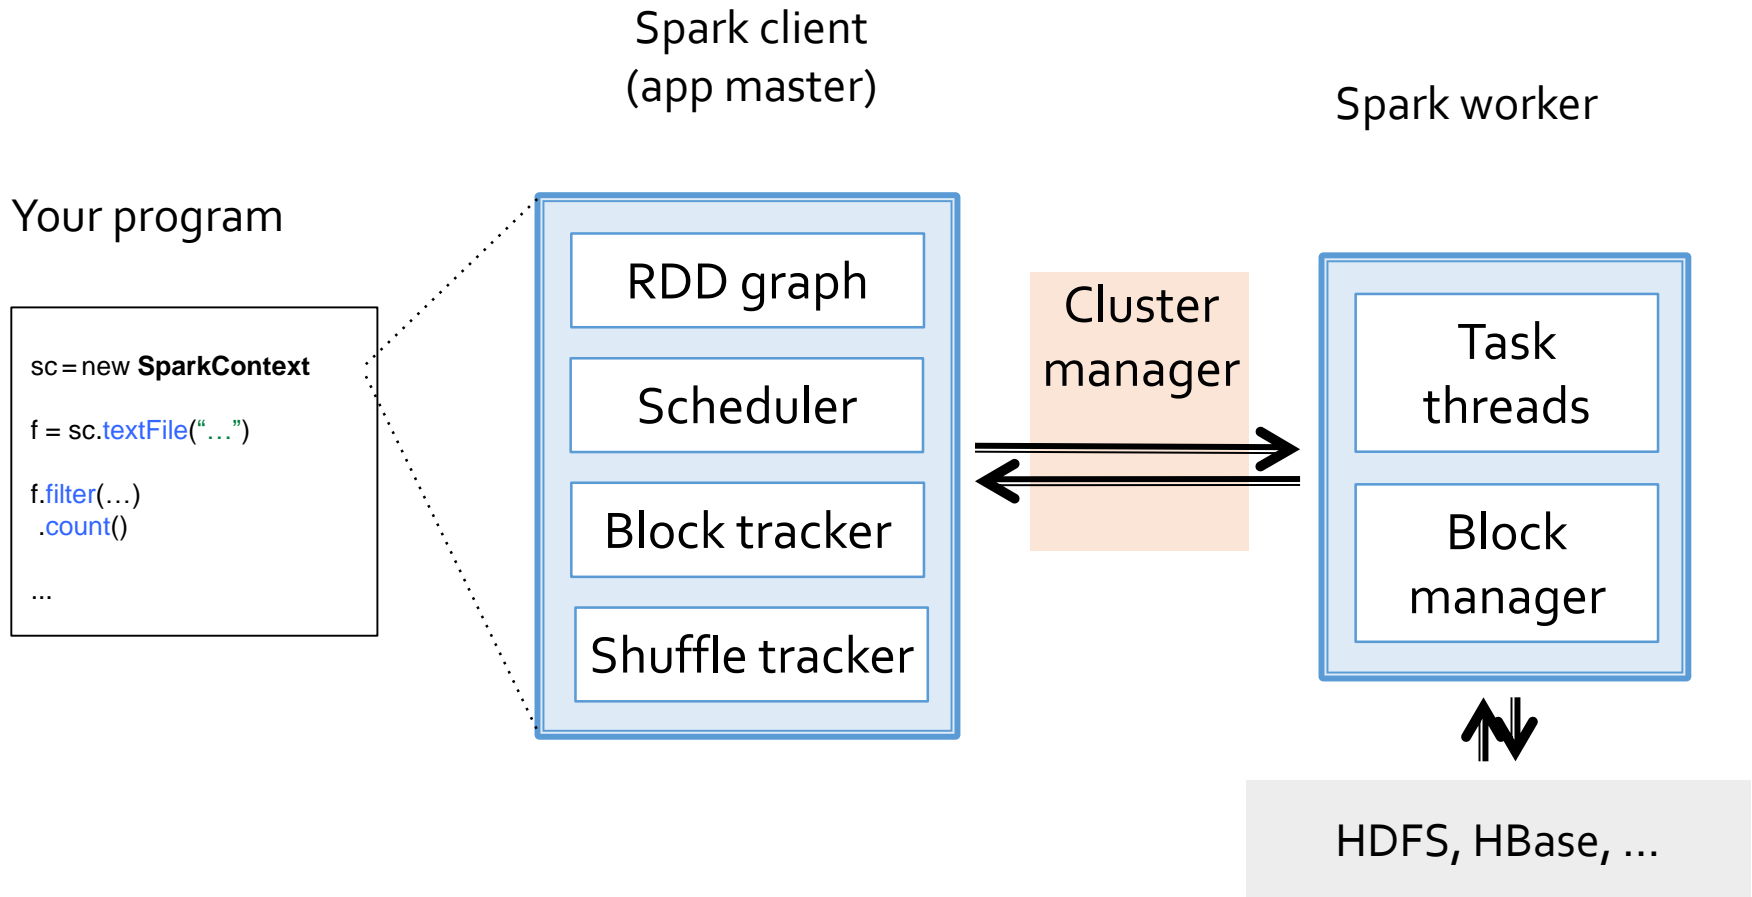
# Software Components

- Spark runs as a library in your program
- Runs tasks locally / on cluster\*
  - Standalone, YARN, Mesos
  - See Cluster Mode Overview\*
- Accesses storage systems via Hadoop API
  - Can use HBase, HDFS, S3, ...



\*=<http://spark.apache.org/docs/latest/cluster-overview.html>

# Components



For more info see video "Introduction to AmpLab Spark Internals" (<https://www.youtube.com/watch?v=49Hr5xZyTEA>) and read slides <http://files.meetup.com/3138542/dev-meetup-dec-2012.pptx>

# Example Job

```
sc = SparkContext (appName="PythonExample")
```

```
file = sc.textFile("hdfs://...")
```

RDDs



```
errors = file.filter(lambda line:"ERROR" in line)
```

```
errors.cache()
```

```
errors.count()
```

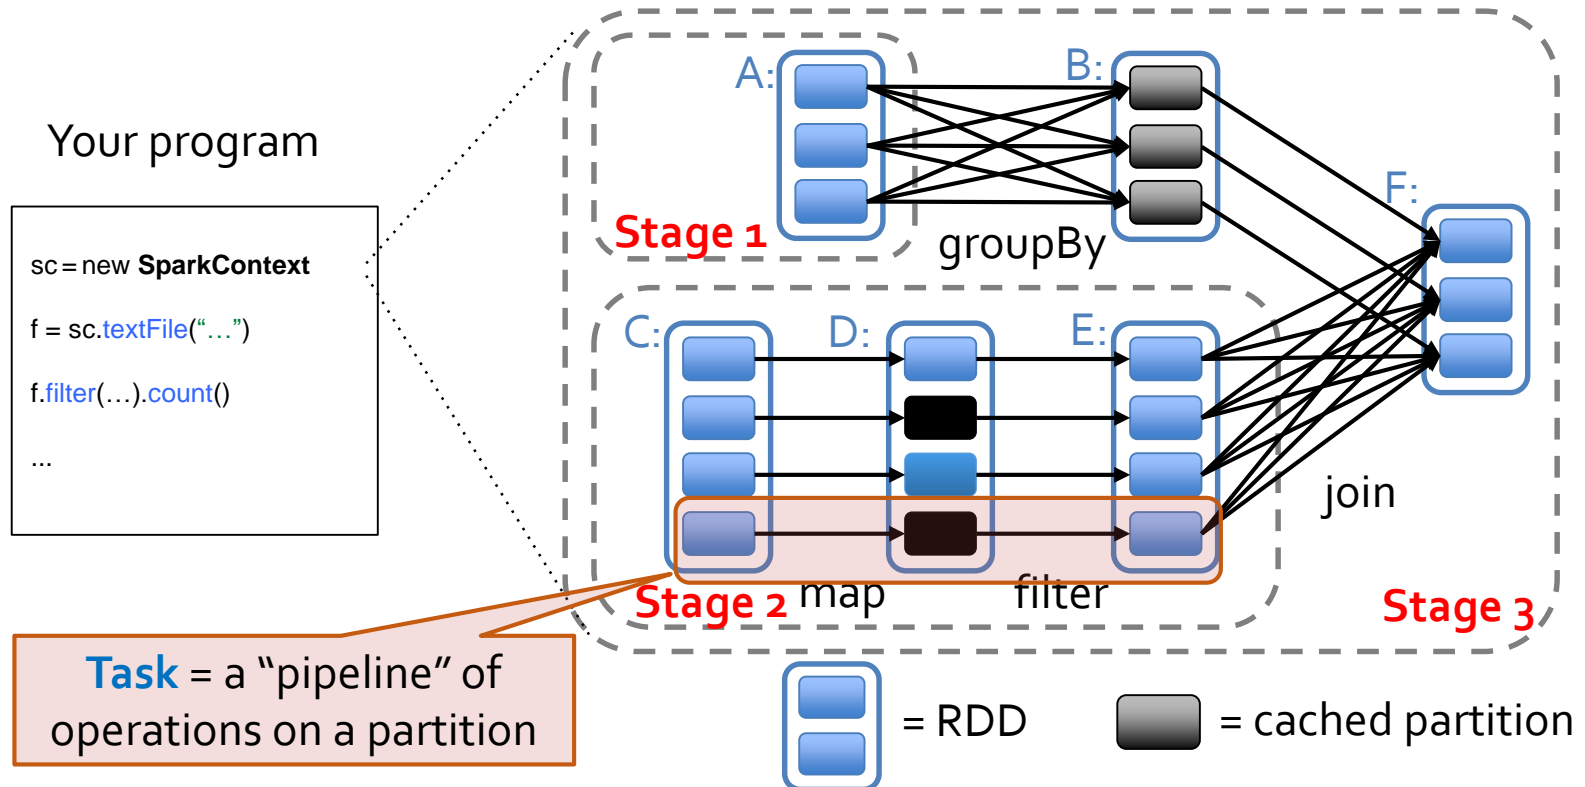
Action



# Operator DAG


The **operator DAG** (Directed Acyclic Graph) captures RDD dependencies

**Stage**: explained later



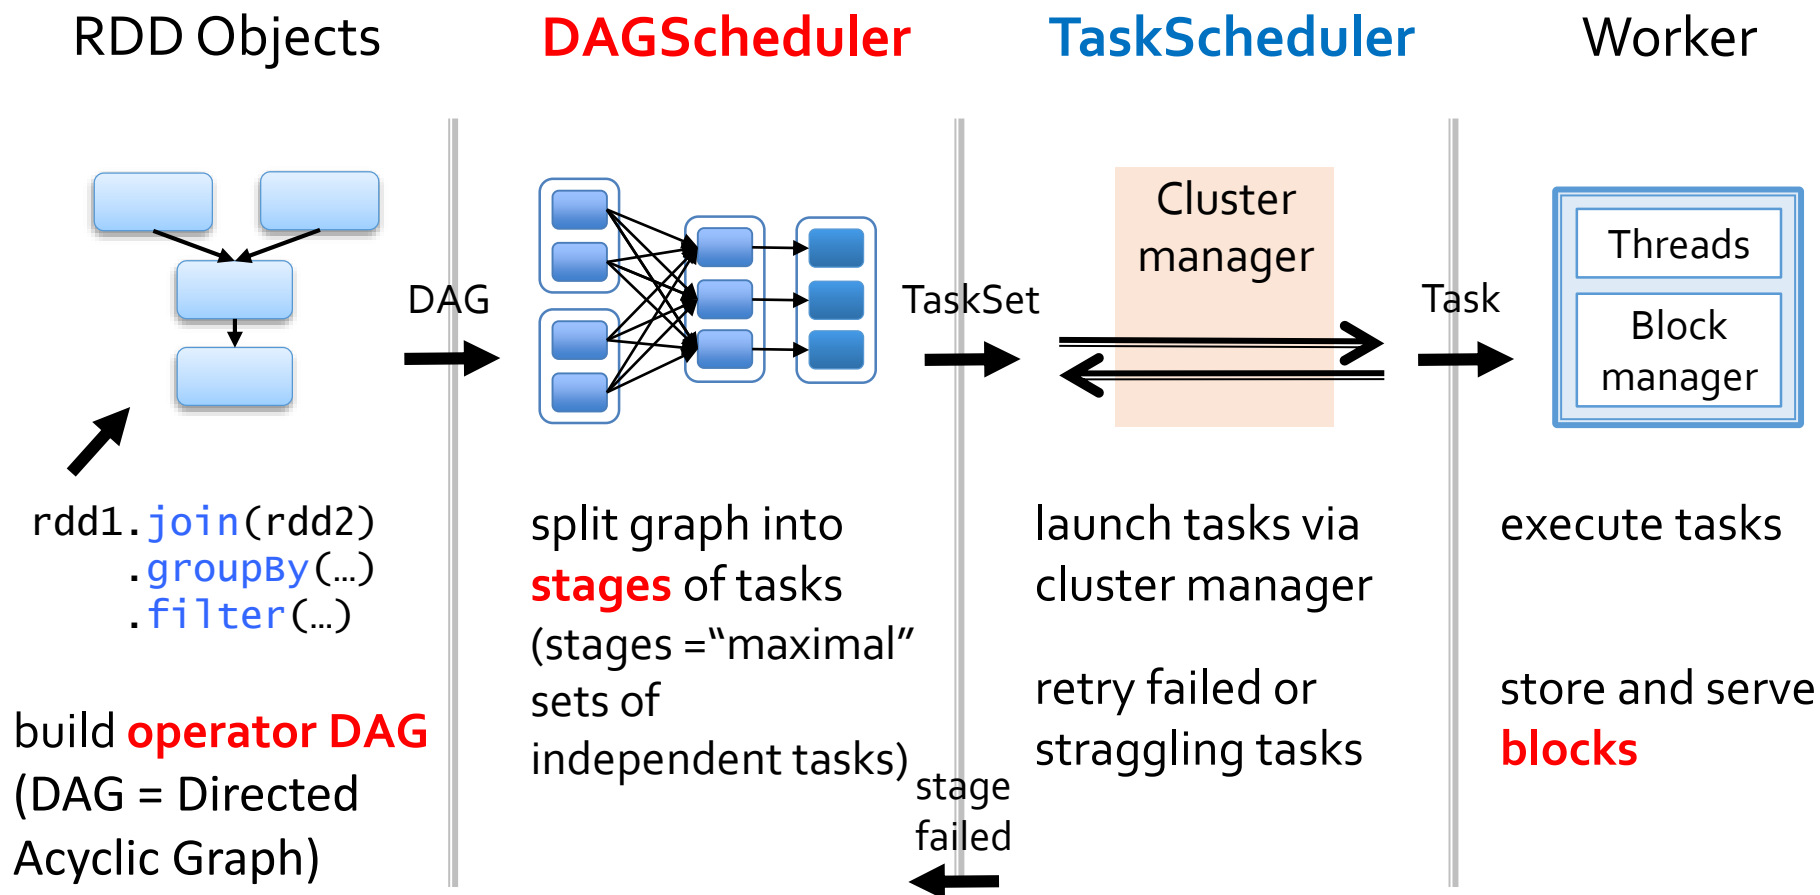
# Stages

- A set of independent tasks, as large as possible
- Stage boundaries are only at input RDDs or “shuffle” operations (like groupBy\*, join, ...)






# Scheduling Process




# Dependency Types in DAG


“Narrow” dependencies:



map, filter




union




join with  
inputs co-  
partitioned

“Wide” (shuffle) deps:



groupByKey



join with inputs not  
co-partitioned

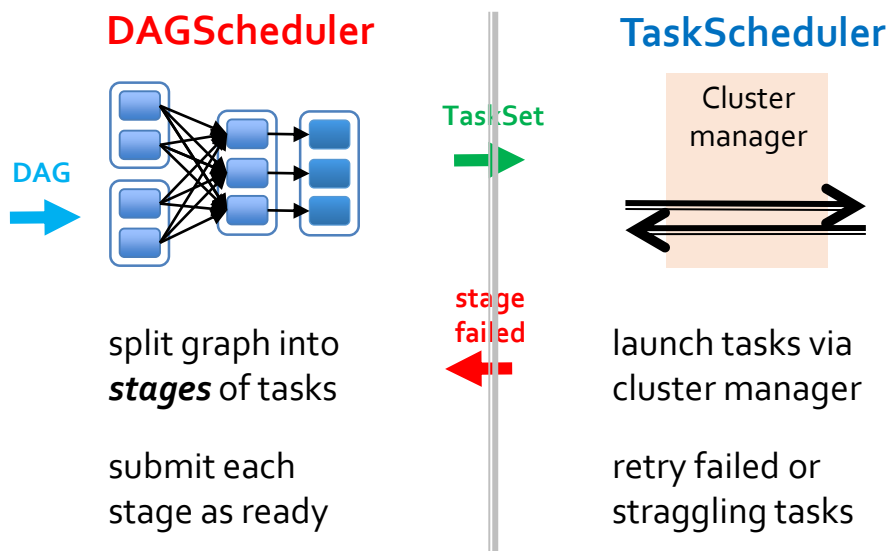
# DAG Scheduler vs. Task Scheduler

## ■ DAG Scheduler – “higher level”

- Builds **stages** of task objects (by code + preferred location)
- Submits them to TaskScheduler as ready
- Resubmits failed stages if outputs are lost

## ■ TaskScheduler


- “Lower level” – similar to Hadoop master
- Given a set of tasks, runs it and reports results
- Exploits data locality
- Local / cluster implementation



# Scheduler Optimizations

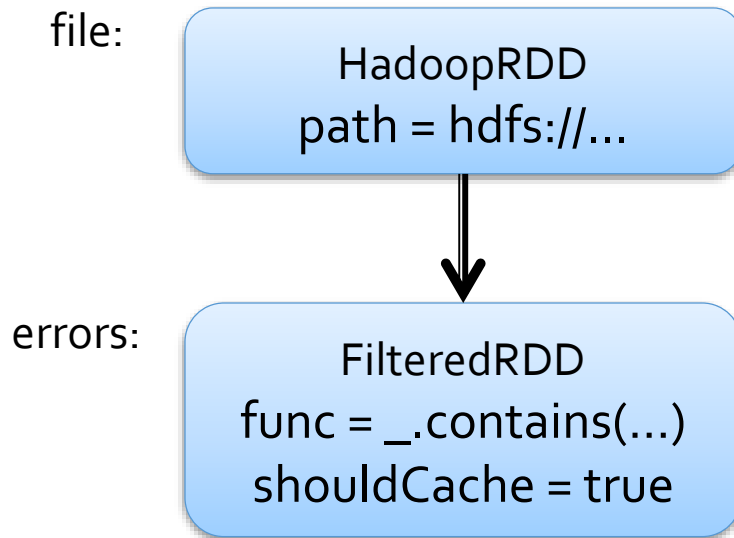
- Pipelines narrow ops. within a stage
- Picks join algorithms based on partitioning (minimize shuffles)
- Reuses previously cached data

In MapReduce, each M-R phase is “individual”  
=> Less optimization!




# RDD Graph

## Dataset-level view:



## Partition-level view:



- **Partition**: a subset of RDD, usually corresponding to a block of HDFS (or other file system)
- **Task**: a “pipeline” of operations on a single partition

# RDD Interface

- Set of partitions (“splits”)
- List of dependencies on parent RDDs
- Function to *compute* a partition given parents
- Optional *preferred locations*
- Optional *partitioning info* (Partitioner)


Captures all current Spark operations!

# Example: HadoopRDD

- `partitions` = one per HDFS block
- `dependencies` = none
- `compute(partition)` = read corresponding block
- `preferredLocations(part)` = HDFS block location
- `partitioner` = none

# Example: JoinedRDD

- `partitions` = one per reduce task
- `dependencies` = “shuffle” on each parent
- `compute(partition)` = read and join shuffled data
- `preferredLocations(part)` = none
- `partitioner` = `HashPartitioner(numTasks)`



Spark will now know  
this data is hashed!



**Thank you.**

Questions?